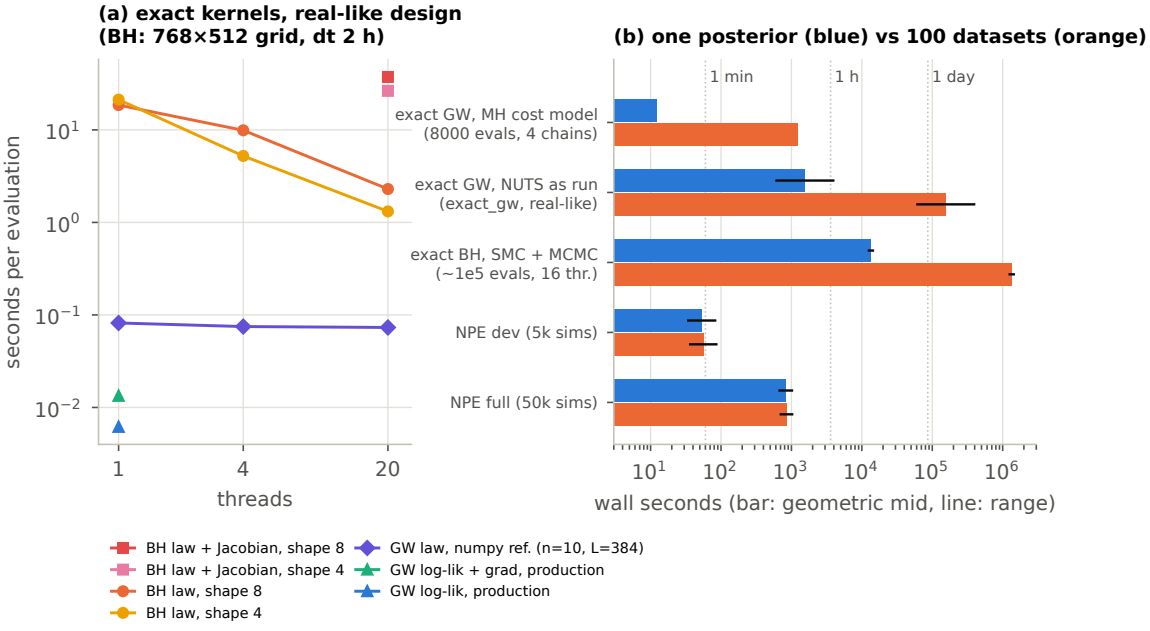
(a) exact kernels, real-like design
(BH: 768×512 grid, dt 2 h)
seconds per evaluation
10^{1}
10^{0}
10^{-1}
10^{-2}
1
4
20
threads
(b) one posterior (blue) vs 100 datasets (orange)
1 min
1 h
1 day
exact GW, MH cost model
(8000 evals, 4 chains)
exact GW, NUTS as run
(exact_gw, real-like)
exact BH, SMC + MCMC
(~1e5 evals, 16 thr.)
NPE dev (5k sims)
NPE full (50k sims)
10^{1}
10^{2}
10^{3}
10^{4}
10^{5}
10^{6}
wall seconds (bar: geometric mid, line: range)
BH law + Jacobian, shape 8
BH law + Jacobian, shape 4
BH law, shape 8
BH law, shape 4
GW law, numpy ref. (n=10, L=384)
GW log-lik + grad, production
GW log-lik, production