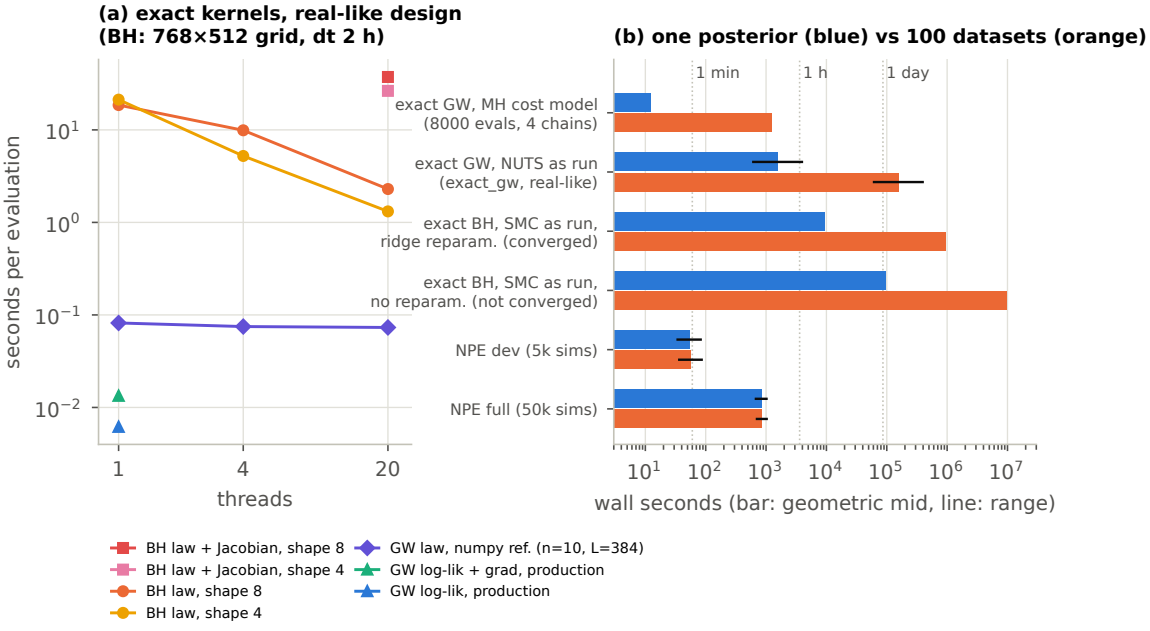
(a) exact kernels, real-like design
(BH: 768×512 grid, dt 2 h)
seconds per evaluation
10^{1}
10^{0}
10^{-1}
10^{-2}
1
4
20
threads
(b) one posterior (blue) vs 100 datasets (orange)
1 min
1 h
1 day
exact GW, MH cost model (8000 evals, 4 chains)
exact GW, NUTS as run (exact_gw, real-like)
exact BH, SMC as run, ridge reparam. (converged)
exact BH, SMC as run, no reparam. (not converged)
NPE dev (5k sims)
NPE full (50k sims)
10^{1} 10^{2} 10^{3} 10^{4} 10^{5} 10^{6} 10^{7}
wall seconds (bar: geometric mid, line: range)
BH law + Jacobian, shape 8
BH law + Jacobian, shape 4
BH law, shape 8
BH law, shape 4
GW law, numpy ref. (n=10, L=384)
GW log-lik + grad, production
GW log-lik, production