BIO-SBI

Simulation experiments for planning organoid lineage experiments

Sequential edition: the experiments in the order of the argument

Two-type branching models ($\mathrm{LGR5}^+ = \mathrm{A}$, $\mathrm{KRT20}^+ = \mathrm{B}$), neural posterior estimation, identifiability and robustness

Status of results	6 November 2026
Repository	`BIO-SBI`, branch `main`, results up to commit `7f4a11f`
Companion	`report/report.pdf` (same content, ordered by experiment ID)
Data	synthetic only (no measured data used)
Audience	group planning the wet-lab organoid experiments

How to read this report. The executive summary (p. 1) is self-contained and follows the same line of argument as the body. Section 1 shows how the experiments depend on each other and maps every experiment ID to its section here and in the ID-ordered report. Section 2 defines the model and every diagnostic. Each experiment section opens with a blue *Why this next* box (the earlier result that raised its question), followed by *Question*, *Setup*, *Result*, *Comment*, and ends with a green *Leads to* box. Orange boxes mark experiments that were planned but never run. Parameters: A→AA etc. are outcome probabilities per cell event; τ_A, τ_B are mean cell lifetimes (Bellman–Harris only).

Executive summary

Scope. All simulation sweeps of E0–E4 are finished (1096 runs, incl. the E4 extensions to 1024–4096 organoids per endpoint and a 20k-simulation check at 4096), all on synthetic data from one reference parameter vector $\theta^\star$ (Table 2), 5 seeds per setting. Two later experiments test the role of the prior with truths swept or drawn from the prior: the E2c rerun (90 runs, §25) and X5 (355 runs; X5a §4, X5c §5, X5d §10, X5b §11). Galton–Watson (GW, time in generations; mapping to days nominal) and Bellman–Harris (BH, time in hours; real-like design = days 6 and 10 = 144/240 h with 165/101 organoids). Every error is compared with the well-specified control of the same sweep.

Key findings, in the order of the argument

I. Identifiability and the prior (Part I).

1. **B's row is prior-dominated if all transitions are free** (E0): B→A = 0.20 vs truth 0.80 at 5k–50k simulations, every design (uniform Dirichlet puts 0.03% mass above 0.8; a flat prior gives 0.53). X5 shows this is structural, not a tail effect: the estimate barely follows the truth (0.04 → 0.22 for truths 0.1 → 0.9) and 1024 organoids do not fix it.
2. **Constraints, not priors or more organoids** (X5a, X5c): with B divisions free no prior (uniform, Jeffreys, flat) and no sample size (up to 1024 organoids) recovers B→A; with them fixed every truth 0.1–0.9 is recovered under every prior. A weak biological prior (A mostly divides, B rarely) is the best soft option (B-division mass 0.54 → 0.11, lowest error, calibrated) but recovers B→A only partly (0.59 at truth 0.9; constraint 0.79).

II. Design and amount of data (Part II).

3. **Endpoints matter more than organoids, in both simulators** (design comparison, E4): days/generations {4,6,10} beat every other set at every size; one late endpoint plateaus; a late extra endpoint ({6,10,14}) adds little. GW: contraction 0.80 → 0.95 → 0.96, MAE 0.052 → 0.023 → 0.019 → 0.020 (16 → 512 → 2048 → 4096 organoids/endpoint); with {4,6,10} every probability except A→AB (error 0.02–0.03, intermittently covered) is recovered by ≈ 128, and at 5k simulations multi-endpoint designs plateau at ≈ 1024 per endpoint (MAE ≈ 0.02; flat to 4096). A single endpoint never catches up: at 4096 organoids MAE 0.046, worse than three endpoints at 32 (0.042). BH: 0.50 → 0.79 → 0.80 → 0.82, MAE 0.114 → 0.060 → 0.058 → 0.057. **The plateau is set by the network, not the data**: at 4096 organoids per endpoint, 20k instead of 5k training simulations on the same datasets lower GW MAE 0.020 → 0.017 and SBC error 0.064 → 0.044, and raise BH run coverage 0.76 → 0.93 (§8.1).
4. **The well-specified GW model is recovered**: at the real-like design MAE 0.033 over the 9 outcome probabilities, contraction 0.89, A→AA 0.52 (truth 0.55); across all E4 designs A→AA 0.47–0.55 — provided B's divisions are fixed to 0 (E0).

III. Bellman–Harris timing (Part III).

5. **BH baseline bias = lifetime/self-renewal ridge.** Well-specified BH gives A→AA ≈ 0.32 and τ_A ≈ 16 h (truth 0.55, 20 h), 0/5 covered, at every design. Fixing lifetimes to the truth gives 0.51 (5/5), MAE 0.109 → 0.059 (§9). With the standard 5k-simulation NPE the posterior stays on the ridge up to 4096 clones per endpoint (up to 12 288 organoids, E4): A→AA 0.29–0.41, τ_A 15.7–17.1 h, 0/5 covered in every design, τ_A contraction 0.99 — confidently on the wrong point. But with 20k simulations at 4096 clones (days 4/6/10) it moves most of the way to the truth: A→AA 0.40 → 0.50 (3/5 covered), τ_A 17.0 → 18.9 h (3/5), lifetime error 0.27 → 0.15 (§8.1). So the data at very large designs can resolve the ridge; at realistic sizes (untested at 20k) lifetime information remains the practical fix.
6. **Averaged over the prior, both simulators are calibrated** (X5d): 90% coverage, SBC ≈ 0.03; contraction 0.68 GW vs 0.51 BH — single-truth failures are about truths and priors, not broken inference. A→AA is learned well on average: the BH ridge bias is specific to this tail truth.
7. **Lifetime information is the practical lever at realistic sizes** (X5b, E3a): a right literature lifetime prior (±25%) moves A→AA 0.32 → 0.46 (MAE 0.109 → 0.068); a lifetime prior 50% too long is detected (support ≈ 0, PPC 0.27), and so is a fixed global lifetime longer than A's; a slightly too short one biases silently. Lifetime CV 0.18–0.5 is harmless, near-exponential cycles are not (E3b); quiescence is not identifiable (E3c).

IV. Misspecified biology and assay (Part IV).

8. **Founder composition**: founder type not learnable in BH (contraction ≤ 0.12), partly in GW (0.68 for truth 0.5); inferring it removes no shift. Founder size: stays at prior in BH; in GW it tracks the truth (0.71 for $p_1 = 0.6$) and inferring it removes the misfit and halves the shifts.
9. **Observed-organoid fraction** (detected/seeded): 0.94–0.95 at $\theta^\star$; dominated by A death (0.3 ⇒ 0.54–

0.58); GW exact and BH simulated agree to 0.02 on average; A and B death trade off (one constraint, not two).

10. **Size of the tested errors (vs paired control; BH | GW):**
    - founders 50% B: A→AB −0.11, τ_B +4 h | shifts ≤ 0.07 (silent in both);
    - 40% multi-cell clones: τ_A −1.8 h, B→A −0.07 | A→AA +0.13, B→A −0.25 (0/5 covered, flagged);
    - B death fixed at 0.4 (true 0.1): B→A −0.19 | −0.23 to −0.26; A death fixed at 0.2 (true 0.03): A persist −0.10, τ_A −2 to −3 h | no fit; deaths fixed at 0: ≤ 0.07 in both (silent);
    - mis-typing ignored, $e = 0.2$: A→AB +0.09 (0.7 sd) | +0.15 (2.1 sd; already 1.4 sd at $e = 0.1$); modelled: information loss (A→AB contraction 0.30 → −0.22 | 0.85 → 0.37);
    - BH only: lifetime CV 0.18–0.5 harmless, exponential A→AA +0.16, B→A −0.18; quiescence (A persist) not identifiable.
11. **Hidden third type (GW, E2c): resolved by the rerun with prior-drawn truths** (§25). The original failure of even the true model (support ≈ 0) was the tail truth (prior probability 10^{-6}–10^{-10}), not the method. In the prior bulk, ignoring C is flagged when C is read as B (9/10 runs at $q = 0$, 7/10 at $q = 0.5$) but *not* when C reads as A (2/10 at $q = 1$, while A/B MAE 0.18 and coverage 0.42 at contraction 0.86: silent confident bias). Modelling C (free or trajectory-constrained) is calibrated (PPC 0, coverage 0.92–0.93) but not identifiable at this design: contraction 0.16–0.20, A/B MAE 0.055–0.059 = prior-mean baseline 0.057; the wrong trajectory constraint is not detected.
12. **A free time scale hides errors** (Figs. 30, 31, §26). In BH the PPC stays ≤ 0.033 for every error except both deaths fixed wrong (0.096) and a global lifetime longer than A's (0.35–0.51). GW, without lifetimes, flags more: A death fixed too high (support ≈ 0), B death fixed at 4× truth (support 0.13–0.33), 40% multi-cell founders (PPC 0.12, support 0.16). Still silent in GW: wrong founder type, B death at 2× truth (B→A 0/5 covered), ignored mis-typing up to 2 posterior sd.

Recommendations for the wet-lab design

- **Endpoints:** three, including an early one — e.g. days 4, 6, 10 (best in both simulators). A single day-10 endpoint is the weakest design, and no number of organoids compensates (GW: 4096 per endpoint at one endpoint is worse than 32 per endpoint at three).
- **Organoids:** ≈128 per endpoint. In GW this recovers every outcome probability with {4,6,10} (A→BB from 32, B persist from 64, B→A from 128; A→AB within 0.03 but only intermittently covered); gains stop at ≈1024 per endpoint. In BH it gives contraction 0.71; A→BB and B death need ≥ 128, A→AB ≥ 256; self-renewal is not recovered at realistic sizes without lifetime information (with 5k simulations not even at 4096 per endpoint; at 4096 with 20k simulations it largely is, §8.1).
- **Measure independently:**
  (a) *A (LGR5$^+$) division time* by live imaging — the single most valuable measurement (it is what separates BH, MAE 0.109, from the GW target, 0.033); per type, not pooled;
  (b) *seeding*: number of organoids seeded per well (observed fraction, sensitive to A death) and whether wells contain multi-cell seeds (≳15% multi-cell clones bias self-renewal and B→A in GW; then fix or infer the founder-size distribution);
  (c) *mis-typing rate* (2D vs 3D comparison) and pass it to the model: ignored, it is tolerable only up to $e \approx 0.05$ once posteriors are as narrow as in GW;
  (d) *cell-cycle variability*: only whether CV is < 0.5 or near-exponential;
  (e) *founder type* and *quiescence* (e.g. Ki67) if they matter biologically — the counts cannot provide them;
  (f) *apoptosis* only if death rates are to be fixed; otherwise keep them free (a value too high biases B→A; in BH silently).
- **Biological constraints help more than data**: B divisions = 0 (if KRT20$^+$ cells do not divide) is what makes B→A identifiable; neither a different prior nor 1024 organoids replaces it (X5a, X5c). A soft biological prior ("B rarely divides") helps but does not replace it; with the constraint keep B's prior near-uniform (an informative symmetric one shrinks B→A toward 1/3).
- **Simulation budget:** for large designs (≳1024 per endpoint) and for any BH fit where the lifetime/self-renewal ridge matters, train on ≥20k simulations with a higher epoch cap (the 20k runs stopped at the 200-epoch cap) and check SBC; 5k results at large sizes are lower bounds on attainable accuracy (§8.1).
- **Literature division times** can enter as an informative lifetime prior (±25% gets most of the benefit of fixing them, X5b); check PPC and support mass for conflict (a prior that is too long is flagged; too-short lifetimes bias silently, E3a).
- **Hidden third type**: endpoint counts of two labels cannot identify a three-type model (E2c rerun). If C may exist, find out how it is labelled: read as B, ignoring it is usually flagged; read as A, it biases the A/B rows silently. Modelling it needs a C-specific marker, trajectory data that justify constraints, or much richer designs — a free C model returns the prior.

Methodological lessons

- **Prior tails**: a truth with $< 2\%$ prior mass is pulled to the prior only where the data are weak or confounded along a ridge (A→AA = 0.55 in BH); where the model is identifiable the data win at every truth position (X5a, B divisions fixed: B→A 0.1–0.9 recovered under every prior). Where it is structurally confounded (B divisions free) the prior decides at every position, and neither a flat prior nor more organoids fixes it. With 10^{-6} or less (original E2c) amortised NPE extrapolates outside the prior support, while the same model is calibrated for truths in the bulk. SBC and prior-averaged coverage (X5d) look fine in all these cases: report prior quantiles of the truth, check prior sensitivity, and compare hand-picked truths with prior-drawn ones.
- **Simulation budget vs data size — confirmed**: at 1024–4096 organoids per endpoint the SBC error of the 5k-simulation amortised NPE doubles (GW 0.03 → 0.06; BH up to 0.039) and accuracy is flat from 2048 to 4096. The 20k check at 4096 (same datasets) separates the two: GW MAE 0.020 → 0.017, SBC 0.064 → 0.044; BH coverage 0.76 → 0.93, A→AA 0.40 → 0.50. The plateau was the network's, and an apparently data-limited ridge can be a training limit (E4, §8.1). Training hit the 200-epoch cap, so 20k is not the limit either.
- **Ridges and free time scales**: endpoint counts fix the growth rate; lifetime, self-renewal and quiescence trade off along it (posterior corr. +0.7, −0.6). The prior picks the point, and the free lifetime also absorbs (hides) errors that GW detects (E1a number, E1c).
- **Compare with paired controls**, not with the truth (BH baseline is biased); 5 seeds resolve shifts of ≈ 0.02–0.03 (probabilities) and 0.5–1.5 h.
- **Diagnostics**: SBC checks the inference model only; PPC and support mass are relative checks and vary between well-specified GW runs (support 0.07–0.89, PPC up to 0.10); support mass ≈ 0 = the model cannot explain the data (numbers from such runs are not estimates).

Planned, never run (Part V): X2 seed spread; X1 methods comparison, X3 GW fit to BH data, X4 combined errors, R real CRC0322 data, P0 BH timing; also asymmetric mis-typing, and the decisive ridge follow-up: 20k simulations with a higher epoch cap at realistic BH sizes (real-like, 512 per endpoint).

Contents

Executive summary 1

1 Roadmap: how the experiments follow from each other 6

2 Methods overview 8

I Identifiability and the role of the prior 10

3 E0 — Identifiability of B's row (GW) 10

4 X5a — Role of the prior I: truth position × prior shape (GW) 11

5 X5c — Role of the prior II: data vs prior at a tail truth (GW) 14

II Design and amount of data 15

6 Design comparison — one vs three endpoints (GW) 15

7 E4 — Amount of data: organoids and endpoints (GW) 16

8 E4 — Amount of data: organoids and endpoints (BH), and comparison with GW 18

III Bellman–Harris timing: the lifetime ridge 22

9 BH ridge — lifetime vs self-renewal 22

10 X5d — Role of the prior III: performance averaged over the prior (GW and BH) 23

11 X5b — Role of the prior IV: informative lifetime prior, right vs wrong (BH) 25

12 E3a — Per-type division time vs one global value (BH) 26

13 E3b — Desynchronisation: lifetime CV misspecified (BH) 27

14 E3c — Quiescence vs division time (BH) 28

IV Misspecified biology and assay 29

15 E1a — Wrong founder type (organoid death as initial conditions; BH and GW) 29

16 E1a — Wrong number of founder cells (BH) 30

17 E1a — Wrong number of founder cells (GW) 31

18 E1b — Death rates and the observed-organoid fraction (BH and GW) 32

19 E1c — Death rates fixed to wrong values (BH) 33

20 E1c — Death rates fixed to wrong values (GW, 4 × 4 grid) 35

21 E2a / E2b — A↔B mis-typing, known vs ignored (BH) 36

22 E2a / E2b — A↔B mis-typing, known vs ignored (GW) **37**

23 E2c — Hidden third type ignored (GW) **39**

24 E2c — Third type modelled: three-arm comparison (GW) **40**

25 E2c rerun — hidden third type, truths from the prior (GW) **41**

26 Synthesis — which errors are detectable (BH and GW) **43**

V Planned experiments 44

27 X2 — Seed spread: dataset vs training **44**

28 Other planned work (X1, X3, X4, R, P0, follow-ups) **44**

A Device and batch-size benchmark **46**

B Reproducibility **46**

C Glossary **48**

1 Roadmap: how the experiments follow from each other

The experiments were planned as one-error-at-a-time sweeps (`TODO.md`), but they were not independent: several of them were launched because an earlier result raised a question. This report follows that chain. The argument runs in five parts.

I. **Identifiability and the role of the prior** (Part I). The first fits (E0) could not recover B's switching back to A. Before anything else could be measured we had to know whether this was the prior or the model, which X5a (truth position, prior shape) and X5c (sample size) settle: it is structural, and fixing B divisions to 0 is the baseline for everything after.

II. **Design and amount of data** (Part II). With an identifiable model, the design comparison and E4 ask which endpoints and how many organoids. In GW everything is recovered; in BH self-renewal is not recovered with the standard 5k-simulation network at any size up to 4096 organoids per endpoint — only a 20k-simulation network at 4096 moves it most of the way.

III. **Bellman–Harris timing** (Part III). That failure is the lifetime/self-renewal ridge. X5d and X5b ask whether it is the truth's position in the prior and whether lifetime information removes it; E3a–E3c test the timing assumptions such information would rest on.

IV. **Misspecified biology and assay** (Part IV). Against well-understood GW and BH baselines, each real-data error is introduced alone: founders (E1a), death (E1b, E1c), mis-typing (E2a/b), a hidden third type (E2c and its rerun, which reuses the tail-truth lesson of Parts I and III). The GW/BH contrast keeps returning: a free time scale absorbs errors. A synthesis section collects detectability.

V. **Planned work** (Part V): seed spread, methods, GW vs BH, combined errors, real data. Engineering (device benchmark) is in Appendix A.

Figure 1 draws the dependencies; Table 1 lists every experiment with the result that motivated it, what it led to, and its section in this report and in the ID-ordered report (`report/report.pdf`).

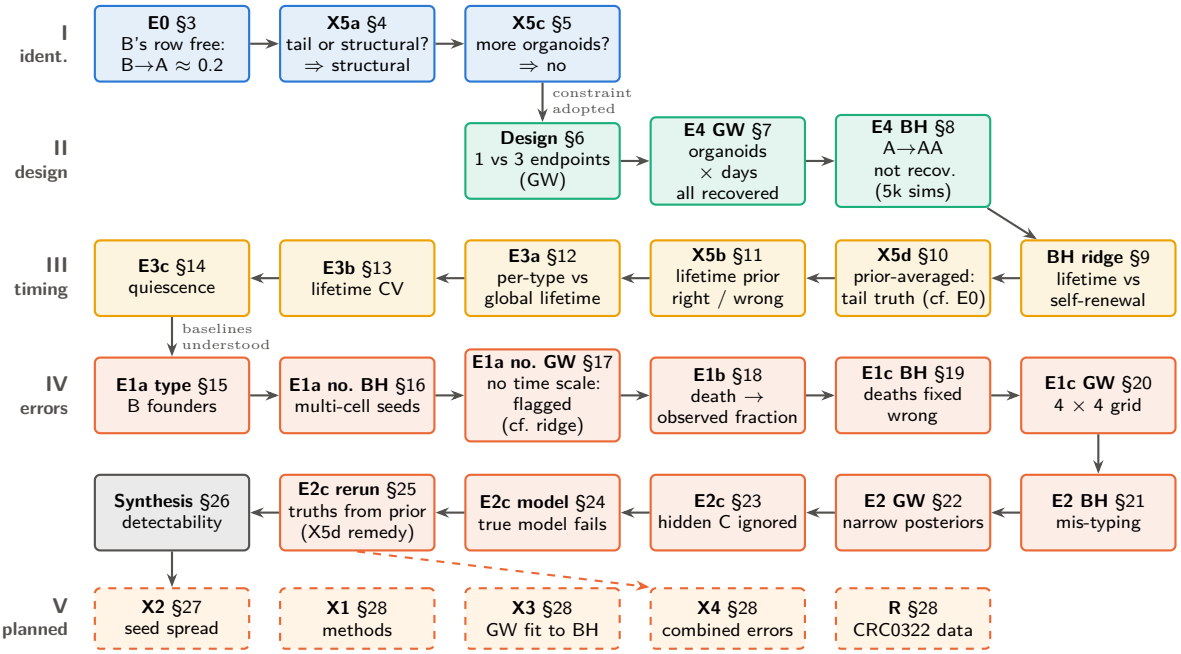


Figure 1: Dependency graph of the experiments. Solid arrows: the order of this report (each result raises the question of the next experiment). Two lessons recur outside the chain and are noted in the boxes: tail truths (E0 → X5d → E2c rerun) and the free time scale (BH ridge → the BH/GW contrasts of Part IV). Dashed arrow: the one-at-a-time errors feed the planned combined-error experiment X4. Colours = parts; dashed orange boxes = planned, never run. Section numbers refer to this report.

Table 1: Experiment index. "Here": section and page in this report. "ID-ordered": section in `report/report.pdf` (same numbers, ordered by experiment ID). Folders under `results/`.

ID	Question (folder)	Here	ID-ord.	Motivated by	Leads to
Part I — identifiability and the role of the prior					
E0	B's row identifiable? (`b_row`)	§3, p. 10	§2	first fits: B→A ≈ 0.2	X5a, X5c; constraint for all later runs
X5a	truth position × prior (`x5/a_*`)	§4, p. 11	§21.1	E0: tail or structural?	X5c; constraint not prior
X5c	organoids vs prior (`x5/c_*`)	§5, p. 14	§21.3	E0 C2ST, X5a	Part II
Part II — design and amount of data					
Design	1 vs 3 endpoints (`design_comparison`)	§6, p. 15	§3	E0/X5a: model fixed	E4 GW
E4 GW	organoids × endpoints (`e4/gw`)	§7, p. 16	§5	design comparison	E4 BH; GW = target
E4 BH	organoids × endpoints (`e4/bh`)	§8, p. 18	§4	E4 GW; BH closer to data	ridge (A→AA not recovered at 5k sims)
E4 20k	5k vs 20k sims at 4096 (`e4/*_xxl_20k`)	§8.1, p. 20	§4.1	E4 plateau, SBC growth	ridge moves; planned 20k at real sizes
Part III — Bellman–Harris timing					
Ridge	lifetime vs self-renewal (`bh_ridge`)	§9, p. 22	§6	E4 BH bias	X5d, X5b; E1a/E1c GW contrasts
X5d	prior-averaged (`x5/d_*`)	§10, p. 23	§21.4	ridge, E0: tail truths	X5b; E2c rerun
X5b	lifetime prior (`x5/b_*`)	§11, p. 25	§21.2	ridge, X5d	E3a
E3a	per-type vs global lifetime (`e3a`)	§12, p. 26	§18	ridge, X5b	E3b
E3b	lifetime CV (`e3b`)	§13, p. 27	§19	E3a: assumed shape	E3c
E3c	quiescence (`e3c`)	§14, p. 28	§20	ridge: other directions	Part IV
Part IV — misspecified biology and assay					
E1a type	founder type (`e1a/founder_type`)	§15, p. 29	§7	baselines set; organoid death	E1a number
E1a no. BH	founder number (`e1a/founder_number_bh`)	§16, p. 30	§8	E1a type	E1a no. GW
E1a no. GW	founder number (`e1a/founder_number_gw`)	§17, p. 31	§9	BH silent: time scale?	E1b
E1b	death → observed fraction (`e1b`)	§18, p. 32	§10	E1a: organoid loss	E1c
E1c BH	deaths fixed wrong (`e1c/bh`)	§19, p. 33	§11	E1b: one constraint	E1c GW
E1c GW	deaths fixed wrong (`e1c/gw`)	§20, p. 35	§12	E1c BH silent	E2
E2a/b BH	mis-typing (`e2/*_bh`)	§21, p. 36	§13	clustering errors	E2 GW
E2a/b GW	mis-typing (`e2/*_gw`)	§22, p. 37	§14	BH: wide posteriors	E2c
E2c	hidden C ignored (`e2c/ignore_c`)	§23, p. 39	§15	E2: C a real type?	E2c model
E2c model	C modelled (`e2c/model_c_*`)	§24, p. 40	§16	E2c: attribution	E2c rerun
E2c rerun	truths from prior (`e2c_prior`)	§25, p. 41	§17	E2c model, X5d	synthesis, X4
—	detectability synthesis	§26, p. 43	summary	all of Part IV	Part V
Part V — planned; Appendix					
X2	seed spread (planned)	§27, p. 44	§23	5 seeds everywhere	—
X1, X3, X4	methods, GW vs BH, combined errors; R, P0 (planned)	§28, p. 44	§24	see section	—
Bench	device / batch (`gpu_bench`)	§A, p. 46	§22	E0 run cost	run settings

2 Methods overview

2.1 Model

- **Two cell types**: A = LGR5$^+$ (stem-like), B = KRT20$^+$ (differentiated).
- **One event per cell**: a cell is replaced by 0, 1 or 2 daughters. Outcomes of A: death, persist (stays, no division), A→B (switch), divisions A→AA, A→AB, A→BB. B mirrors A.
- **Parameters**: one probability simplex per type; Dirichlet(1,...,1) prior (uniform).
- **Laws**: `two_type` = all 12 outcomes free (10 free params); `reversible_ab` = B divisions fixed to 0 (B: death, persist, B→A; 7 free) — the baseline of every experiment after E0.
- **Founders**: a clone (organoid) starts from N cells ($P(N{=}n)$ = `founder_size`) of i.i.d. type (`founder`); baseline: one A cell.

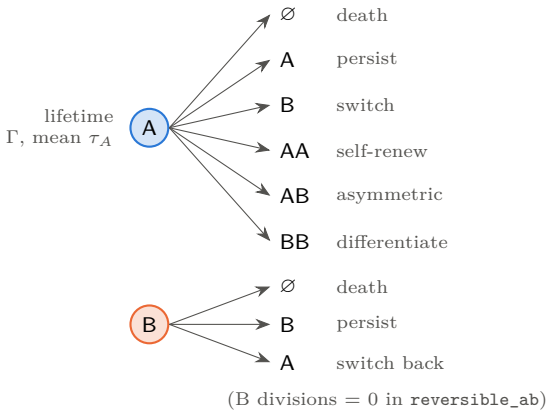


Table 2: Reference truth $\theta^\star$ used by every experiment unless stated (`law.reference`); one founder cell of type A. Prior probability of exceeding the truth under the uniform Dirichlet: A→AA 1.8%; B→A 0.03% (`two_type`) or 4% (`reversible_ab`).

Type	death	persist	switch	→AA	→AB	→BB	mean lifetime (BH)
A	0.03	0.03	0.05 (A→B)	**0.55**	0.05	0.29	τ_A = 20 h
B	0.10	0.10	**0.80** (B→A)	0	0	0	τ_B = 30 h

BH lifetimes: Gamma, shape 8 (CV 0.35). GW: one event per generation.

2.2 Two simulators

	Galton–Watson (GW)	**Bellman–Harris (BH)**
Time	synchronous generations; endpoints in *generations* (real-like: 6 and 10)	continuous; Gamma lifetimes per type; endpoints in *hours* (real-like: 144 and 240 h = days 6, 10)
Extra parameters	none	τ_A, τ_B (log-uniform prior); shape fixed (8)
Sampling	exact detection-conditioned sampler; exact observed fraction	forward simulation + rejection; observed fraction estimated
Cost	cheap	slower (rejection sampling; P0 timing not run)

GW and BH share the outcome law but not the time unit; a GW generation is *not* a day (X3 needs a mapping).

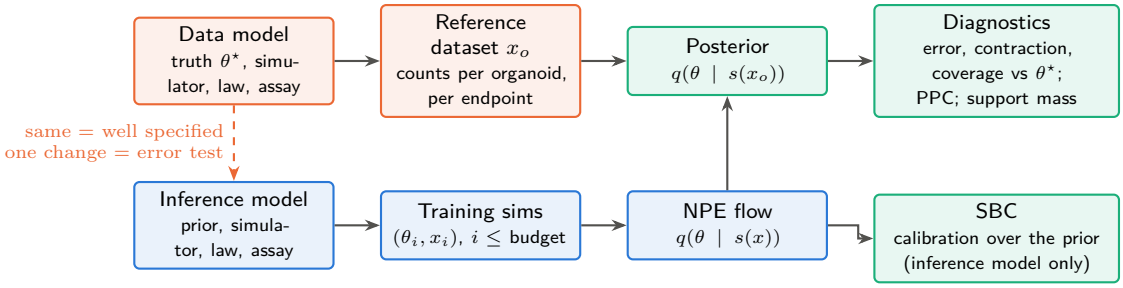
2.3 Observation, data model and inference model

- **Observation**: biology → capture (Binomial per true type) → labelling (confusion matrix true type × label; may be rectangular, e.g. hidden C read as A with prob. q) → selection of *recovered* clones (≥ 1 detected cell). Baseline: perfect capture and typing.
- **Data**: per endpoint, A and B counts of each recovered organoid. Endpoints are destructive: different organoids per endpoint.
- **Misspecification**: the reference data come from a separate *data model* (`data.simulator`, `data.law`, `data.observation`, with seed `data_seed`); the *inference model* defines prior, training simulations and checks. An error is tested by changing one side only.

2.4 Inference

- **NPE** (`sbi` 0.27, neural spline flow, 5 transforms, width 64), amortised, 1 round; sequential NPE (SNPE) only in E0.
- **Summaries** (`moments`): per endpoint, mean, sd, quantiles of $\log(1+n)$ for A, B, total; fraction of clones containing each type; fraction empty; correlation of log counts.
- **Budgets**: `dev` = 5 000 simulations (all BH and most GW error experiments); 20 000 for E0/design comparison; E0 also 50 000. E0 showed 5k→50k does not change recovery. Per run: 2 000 posterior samples, SBC on 200 prior draws, PPC with 200 datasets.
- **Parameterisation**: stick-breaking of each simplex (exact Dirichlet reparameterisation).

2.5 Pipeline and experiment logic



Experiment logic: vary one error on the data side (or one assumption on the inference side), keep everything else at the baseline, and compare with the *well-specified control of the same sweep* (same seeds, same datasets; often the same trained network, so differences are paired by seed).

2.6 Diagnostics and how to read them

Metric	Definition	How to read
Posterior mean / error	mean of 2 000 posterior samples; |mean − truth|. MAE over probabilities only (`mae_probabilities`); lifetimes separately	compare misspecified vs control; "shift" = difference of posterior means, paired by seed, ± paired SE over 5 seeds
Contraction	$1 - \mathrm{Var}_{\mathrm{post}}/\mathrm{Var}_{\mathrm{prior}}$	0 = data uninformative; < 0 = wider than prior; > 0.5 = informative
90% coverage across seeds	number of seeds (of 5) whose 90% credible interval contains the truth	4–5/5 expected; 0/5 with high contraction = confident bias
SBC coverage error	mean |empirical − nominal| coverage of central intervals (50–95%) over 200 prior draws	≤ 0.03 = calibrated. Averages over the prior: blind to tail truths and to misspecification
PPC fraction outside 95%	share of summary features whose observed value lies outside the 95% posterior predictive interval	well specified: 0–0.03. Higher = the model cannot reproduce the data. Relative check
Posterior support mass	share of the NPE posterior inside the prior support	well specified $\approx$0.5–0.9 (not 1). ≈ 0 = model cannot explain the data; few/no samples, not an estimate
Recovered (TODO criterion)	contraction > 0.5, |error| < 0.05, coverage $\geq 4/5$	used for "organoids needed" (E4)

Reading rules used throughout. (1) *BH baseline bias*: well-specified BH runs give A→AA ≈ 0.32 and $\tau_A \approx$ 16 h (truth 0.55, 20 h); BH misspecified arms are compared with their control, never with the truth. (2) With all transitions free, B's row is prior-dominated (E0). (3) A misspecification that raises neither PPC nor lowers support mass is *undetectable from the fit*. (4) 5 seeds per setting: effects below $\approx$2 paired SE are not resolved. (5) Wall times in run metrics were measured under heavy contention; only the device benchmark (§A) measures speed.

Part I

Identifiability and the role of the prior

Motivation. Every later question (how many organoids, which errors matter) presupposes that the parameters can be learned from endpoint counts at all. The very first fits failed on one parameter, B→A. This part establishes why: whether the prior or the model is at fault decides what the remedy is (a different prior, more data, or a biological constraint). It ends with the constraint that all later experiments adopt (B divisions = 0, `reversible_ab`).

3 E0 — Identifiability of B's row (GW)

Why this next. This is the starting point. With the general two-type law (all 12 outcomes free) the first well-specified fits returned B→A ≈ 0.2 for a truth of 0.8. Nothing downstream can be interpreted until this is understood.

Question

With all transitions free, B→A (true 0.80) came out at ≈0.2. Is this under-training, one bad dataset, miscalibration, a prior-tail effect, or non-identifiability (switching vs division)?

Setup

Simulator	Galton–Watson, perfect observation
Law (free)	`two_type` (10) vs `reversible_ab` (7; B divisions fixed to their true 0)
Data vs inference	well specified (same $\theta^\star$, Table 2)
Design	generations 4, 8, 12 × 32 clones (96 organoids)
Inference / budget	NPE amortised 5k, 20k, 50k; SNPE 4 rounds × 5k; `two_type` with flat prior on B→A (Beta(1,1)), 20k
Seeds / runs	5 / 40 (+ direct checks: C2ST of mean-preserving alternatives, 500 datasets per side)
Varied	law, budget, method, prior
Folder	`results/b_row/` (`distinguish.csv`, `truth_vs_posterior_20k*.csv`)

Result

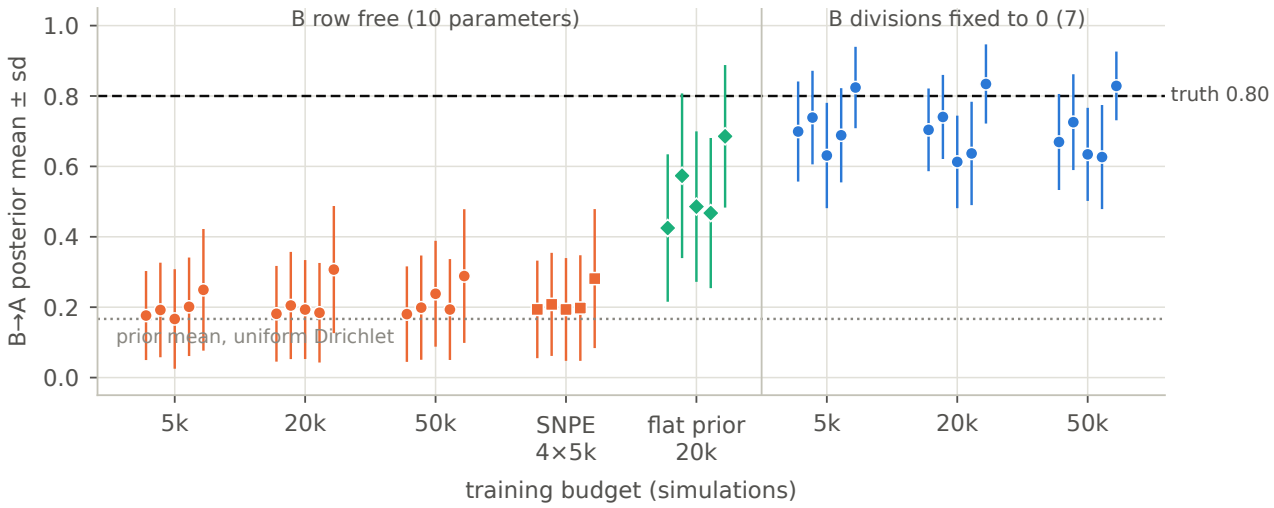


Figure 2: B→A posterior mean ± sd per seed. Orange: all B outcomes free (uniform Dirichlet; squares: SNPE). Green: flat prior on B→A. Blue: B divisions fixed to 0.

Law	Training	B→A mean	post. sd	contraction	covered	MAE (all)
`two_type`	5k / 20k / 50k	0.20 / 0.21 / 0.22	0.14–0.15	−0.03 / −0.13 / −0.18	0/15	0.14
`two_type`	SNPE 4×5k	0.21	0.16	−0.22	0/5	0.14
`two_type`, flat prior	20k	0.53	0.21	0.46	4/5	0.10
`reversible_ab`	5k / 20k / 50k	0.72 / 0.71 / 0.70	0.13–0.14	0.67 / 0.71 / 0.69	15/15	0.05

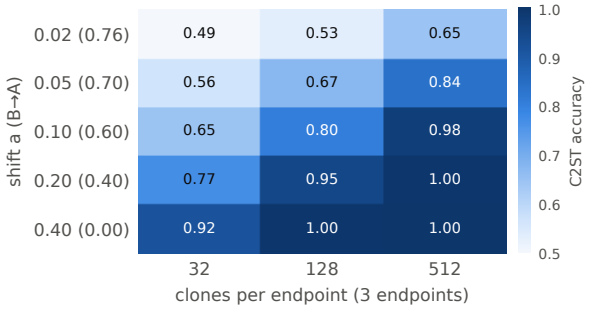


Figure 3: C2ST accuracy separating the truth from a mean-preserving alternative (B→AA = a, B→A = $0.8 - 2a$, B death = $0.1 + a$). 0.5 = indistinguishable.

- SBC coverage error 0.02–0.04 in every arm (looks calibrated).
- PPC ≤ 0.014 in every arm.
- Flat prior: corr(B→A, B→AA) = −0.67; SBC fails (0.069) for B's other coordinates (concentration 0.2), not for B→A.
- Truth-vs-posterior C2ST: inconclusive (0.62–0.95 in both laws).

Comment

- Under-training, dataset and calibration are ruled out: 10× more simulations, 5 datasets, SBC all unchanged.
- Main cause = prior tail: B→A ∼ Beta(1,5), P(> 0.8) = 0.03% (≈ 15 of 50k training simulations). SNPE cannot help: round 0 is already prior-like. *Revised by X5 (§4, §5)*: with B divisions free B→A is not recovered at any truth ≥ 0.5 (bulk included) nor with 1024 organoids, so the confounding below is the root cause and the tail only makes it visible.
- Data weakly confound switching with division: even with a flat prior the posterior is wide (sd 0.21) and anti-correlated with B→AA. They differ only in clone-size variability; 128–512 clones/endpoint make moderate differences detectable.
- Consequence: fix what biology allows (B divisions = 0 ⇒ B→A ≈ 0.70, contraction 0.7); otherwise report prior sensitivity. All later experiments use `reversible_ab`.

Leads to. Two explanations remain open: the prior tail (B→A = 0.8 has 0.03% prior mass) and structural confounding of switching with division (the flat-prior posterior is wide and anti-correlated with B→AA). X5a (§4) separates them by moving the truth through the prior; X5c (§5) asks whether more organoids, which the C2ST suggests would help, can overrule the prior.

4 X5a — Role of the prior I: truth position × prior shape (GW)

Why this next. E0 (§3) attributed the B-row failure mainly to the prior tail, with weak confounding as a second cause. If the tail were the cause, a truth in the prior bulk or a flatter prior should be recovered; if the confounding were the cause, no truth position or prior should help unless B divisions are fixed.

About X5. X5 ("role of the prior") is one experiment in four parts, placed here where each is motivated: X5a (this section), X5c (§5), X5d (§10), X5b (§11). Common to all: synthetic data, real-like design (GW generations 6/10 or BH 144/240 h; 165/101 clones unless swept), perfect observation, one A founder, NPE 5k (`budget=dev`), SBC 200, PPC 200; 355 runs in total (X5a 250, X5b 15, X5c 30, X5d 60). Folder `results/x5/` (overview `info.md`, per-part tables `x5*_findings.md`).

Question

Does recovery of B→A depend on where the truth sits in the prior (0.1–0.9) and on the prior's shape, with B divisions free vs fixed? Is E0's failure a tail-truth effect?

Setup

Data model	GW; B→A ∈ {0.1, 0.3, 0.5, 0.7, 0.9}, B death = B persist = (1 − B → A)/2; other parameters $\theta^\star$
Laws (data = inference)	`two_type` (B divisions free, 10 free) and `reversible_ab` (B divisions fixed to 0, 7 free). Prior P(B→A > 0.9): 10^{-5} resp. 0.01
Priors (inference only)	uniform Dirichlet; Jeffreys (all concentrations 0.5); flat marginal on B→A (Beta(1,1): concentration 0.2 on B's 5 other outcomes in `two_type`, 0.5 on death/persist in `reversible_ab`). *Biological* (Dirichlet pseudo-counts per type, each row sums to 1; total ≈ 5–6 = weakly informative): **bio-B** = B row only, death/persist/to_A 1.5 and B divisions 0.2 each (`two_type`; fixed in `reversible_ab`), A row uniform; **bio** = bio-B plus A row death/persist/to_B 0.5, div_AA 2, div_AB 1, div_BB 1.5. Rationale: encode the biology without the answer — A (stem-like) mostly divides, self-renewal most likely; B (differentiated) rarely divides but may die, persist or switch back. bio-B vs bio separates B's prior from A's
Seeds / runs	5 / 125 per law (250: 150 reference + 100 biological)
Folder	`results/x5/a_two_type/`, `results/x5/a_reversible_ab/`, biological priors in `a_two_type_bio/`, `a_reversible_ab_bio/` (pooled `x5a_groups.csv`, `x5a_runs.csv` in the first two)

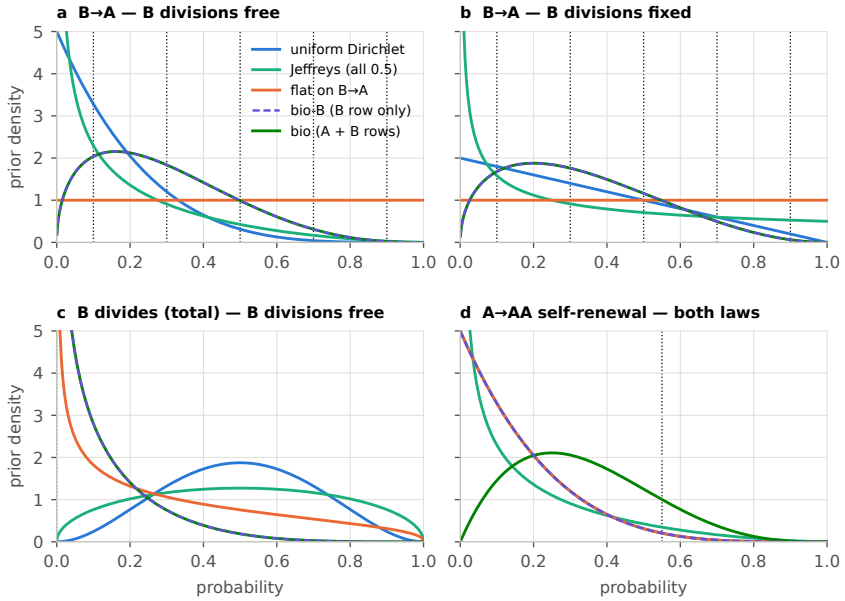


Figure 4: The five X5a priors compared (inference model; exact Beta marginals of the Dirichlet priors). (a, b) B→A with B divisions free / fixed; (c) total probability that B divides (B→AA + AB + BB), B divisions free; (d) A→AA, identical for both laws. Curves that coincide are drawn on top of each other: bio-B (dashed) = bio in (a–c); uniform = flat B→A = bio-B in (d). Dotted: truths used in X5a (B→A 0.1–0.9; B divisions 0; A→AA 0.55). Densities clipped at 5. Beta parameters, quantiles and tail probabilities: `tables/x5a_priors.csv`.

Result

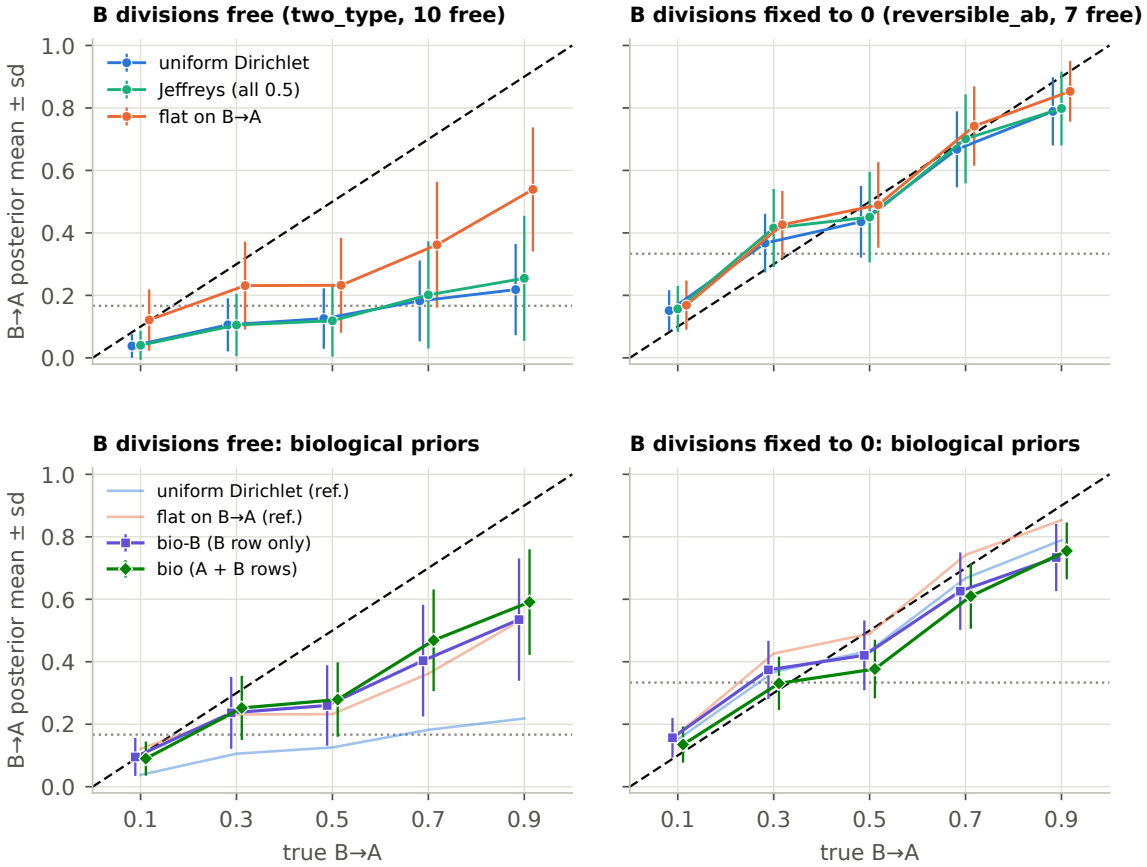


Figure 5: X5a: B→A posterior mean (mean over 5 seeds) ± mean posterior sd against the true B→A, one line per prior (slightly offset). Top: reference priors; bottom: biological priors bio-B and bio, with the uniform and flat-B→A means as thin reference lines. Dashed: identity (perfect recovery). Dotted: uniform-prior mean (1/6 left, 1/3 right).

	B divisions free (`two_type`)			B divisions fixed (`reversible_ab`)		
Truth	uniform	Jeffreys	flat B→A	uniform	Jeffreys	flat B→A
0.1	0.04 (3/5)	0.04 (5/5)	0.12 (5/5)	0.15 (5/5)	0.16 (4/5)	0.17 (5/5)
0.3	0.11 (1/5)	0.11 (2/5)	0.23 (5/5)	0.37 (5/5)	0.42 (5/5)	0.43 (4/5)
0.5	0.13 (0/5)	0.12 (0/5)	0.23 (2/5)	0.44 (5/5)	0.45 (5/5)	0.49 (5/5)
0.7	0.18 (0/5)	0.20 (0/5)	0.36 (3/5)	0.67 (5/5)	0.70 (5/5)	0.74 (5/5)
0.9	0.22 (0/5)	0.26 (0/5)	0.54 (1/5)	0.79 (5/5)	0.80 (5/5)	0.85 (5/5)
SBC error	0.037	0.054	0.080	0.032	0.056	0.037

	B divisions free (`two_type`)		B divisions fixed (`reversible_ab`)	
Truth	bio-B	bio (A+B)	bio-B	bio (A+B)
0.1	0.10 (5/5)	0.09 (5/5)	0.16 (4/5)	0.14 (5/5)
0.3	0.24 (5/5)	0.25 (5/5)	0.37 (4/5)	0.33 (5/5)
0.5	0.26 (1/5)	0.28 (1/5)	0.42 (5/5)	0.38 (4/5)
0.7	0.40 (2/5)	0.47 (3/5)	0.63 (5/5)	0.61 (4/5)
0.9	0.54 (1/5)	0.59 (1/5)	0.73 (1/5)	0.75 (3/5)
run MAE	0.05–0.09	0.03–0.07	0.03–0.06	0.03–0.05
SBC error	0.052	0.055	0.033	0.041
B-division mass	0.21	0.11	—	—

B→A posterior mean (seeds covered by the 90% interval). Run MAE = range over truths of the mean per-run MAE (all free parameters); B-division mass = posterior total of B→AA/AB/BB (truth 0; 0.54 under the uniform and flat priors). Reference priors, run MAE: free 0.12–0.15 (uniform), 0.10–0.16 (flat); fixed 0.02–0.06 (uniform). Numbers: `tables/x5a.csv`.

Comment

- **B divisions free: the estimate hardly follows the truth** (0.04 → 0.22 as the truth goes 0.1 → 0.9, uniform prior); no truth ≥ 0.5 is covered under the uniform or Jeffreys prior, and B→A contraction falls to ≤ 0 at 0.9. The data cannot separate switching from division, so the prior sets the answer — at *every* truth position, not only in the tail.
- A flat prior on B→A tracks better (0.12 → 0.54) but still shrinks heavily and costs calibration (SBC 0.080) and coverage of the other parameters (run coverage 0.62–0.87).
- **B divisions fixed: every truth is recovered under every prior** (B→A contraction 0.73–0.94; the three priors differ by ≤ 0.06, mild shrinkage at the extremes). The prior only matters once the model is structurally confounded.
- Conclusion for E0 (§3): the B-row failure is structural (switching vs division), not a bad truth position; a prior cannot substitute for the constraint.
- **Biological priors.** *B divisions free*: the soft "B rarely divides" prior cuts the posterior B-division mass from 0.54 to 0.21 (bio-B) and 0.11 (bio); B→A tracks like the flat prior (0.54–0.59 at truth 0.9) but with better calibration (SBC 0.05 vs 0.08) and the lowest error of any free-B prior (run MAE 0.03–0.07 for bio; the A-row prior gives most of this gain). It is still far from the hard constraint (0.79, 5/5 at truth 0.9): truths ≥ 0.5 are mostly not covered. *B divisions fixed*: the symmetric 1.5 prior on B's three outcomes (mean 1/3) shrinks B→A toward 1/3, and truth 0.9 is covered in only 1/5 (bio-B) and 3/5 (bio) seeds; the A-row prior lowers the error and raises support mass (0.79–0.89 vs 0.64–0.79 uniform). A soft biological prior helps but does not replace the constraint; with the constraint, keep B's prior near-uniform.

Leads to. The constraint, not the prior, decides. Remaining question: is it only a matter of sample size? (X5c, §5).

5 X5c — Role of the prior II: data vs prior at a tail truth (GW)

Why this next. X5a (§4) showed the confounding at the real-like size (165/101 clones). E0's C2ST (§3) predicted that 128–512 clones per endpoint make moderate switching/division differences detectable, so a large enough experiment might overrule the prior without a constraint.

Question

With all transitions free and a tail truth (B→A = 0.9, prior P = 10^{-5} under Beta(1,5)), how many organoids overrule the prior?

Setup

Data model	GW, `two_type`; B: death 0.05, persist 0.05, B→A 0.9 (B divisions 0); A row $\theta^\star$
Inference	`two_type` (10 free), uniform Dirichlet vs flat marginal on B→A (concentration 0.2 on B's other outcomes)
Design	generations 6/10 × {32, 128, 512} clones per endpoint
Seeds / runs	5 / 30
Folder	`results/x5/c_data_vs_prior/` (`x5c_groups.csv`, `x5c_runs.csv`)

Result

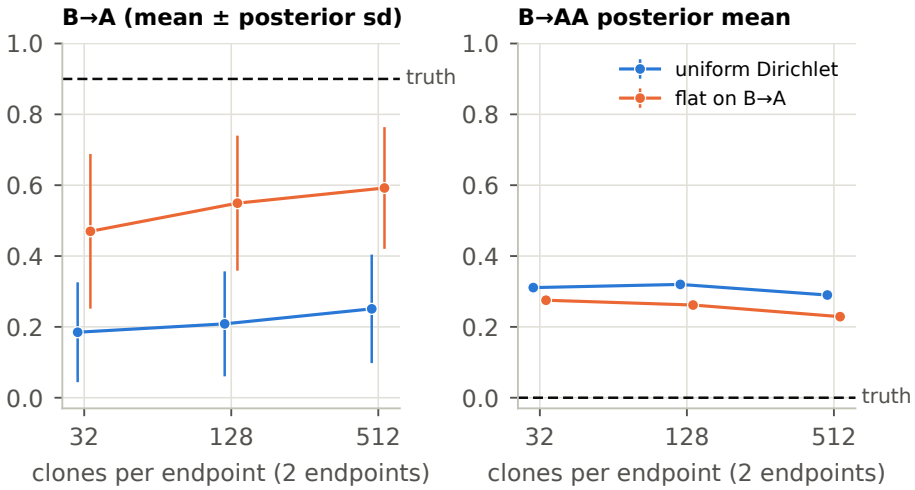


Figure 6: X5c: B→A (mean over 5 seeds ± mean posterior sd) and B→AA posterior mean against clones per endpoint; dashed = truth.

Prior	clones	B→A (0.9)	post. sd	cov.	contr.	B→AA (0)	MAE	support	SBC
uniform	32	0.19	0.14	0/5	0.01	0.31	0.161	0.79	0.031
	128	0.21	0.15	0/5	−0.10	0.32	0.151	0.72	0.037
	512	0.25	0.15	0/5	−0.17	0.29	0.139	0.55	0.044
flat B→A	32	0.47	0.22	2/5	0.44	0.28	0.119	0.49	0.078
	128	0.55	0.19	1/5	0.57	0.26	0.101	0.42	0.081
	512	0.59	0.17	1/5	0.65	0.23	0.086	0.46	0.083

PPC ≤ 0.03 in every group. Numbers: `tables/x5c.csv`.

Comment

- **1024 organoids do not overrule the uniform prior**: B→A moves 0.19 → 0.25 (truth 0.9), never covered, contraction ≤ 0; the mass goes to B→AA ($\approx$0.3, truth 0) — the switching/division confounding of E0.
- A flat prior helps and improves with data (0.47 → 0.59, contraction up to 0.65) but is still far from 0.9 at 1024 organoids and costs calibration (SBC 0.08).
- Endpoint counts of this kind cannot resolve the confounding by sample size at realistic scales (with NPE 5k; E4 later shows a larger simulation budget can matter at very large designs, §8.1); the structural constraint (B divisions fixed, X5a) is what works.

Leads to. Part II: with B divisions fixed (`reversible_ab`) as the baseline, which design and how many organoids? The design comparison (§6) also re-checks the free law at up to 384 organoids.

Part II

Design and amount of data

Motivation. With an identifiable model (Part I), the planning question is the design: which days to measure and how many organoids per day. E0 showed that the information lies in clone-size variability across time, which suggests that the number of endpoints may matter as much as the number of organoids. We first compare one and three endpoints in GW, then map the full grid in GW and BH. The BH half produces the result that drives Part III.

6 Design comparison — one vs three endpoints (GW)

Why this next. Part I fixed the law (B divisions = 0, §4). E0 used 3 endpoints × 32 clones; the real data have 2 endpoints, and a wet-lab plan can trade endpoints for organoids. Which matters more, at equal total organoids?

Question

Which planned design identifies the model best, and is it the number of time points or the number of organoids that matters?

Setup

Simulator	Galton–Watson, perfect observation
Law (free)	`two_type` (10) and `reversible_ab` (7)
Data vs inference	well specified, $\theta^\star$
Design	`single_128`: gen. 12 × 128; `three_128`: gen. 4, 8, 12 × 128; `single_384`: gen. 12 × 384; reference `three_32` (from E0)
Inference / budget	NPE amortised, 20k
Seeds / runs	5 / 30 (+ 10 reference runs from E0)
Varied	design × law
Folder	`results/design_comparison/`

Result

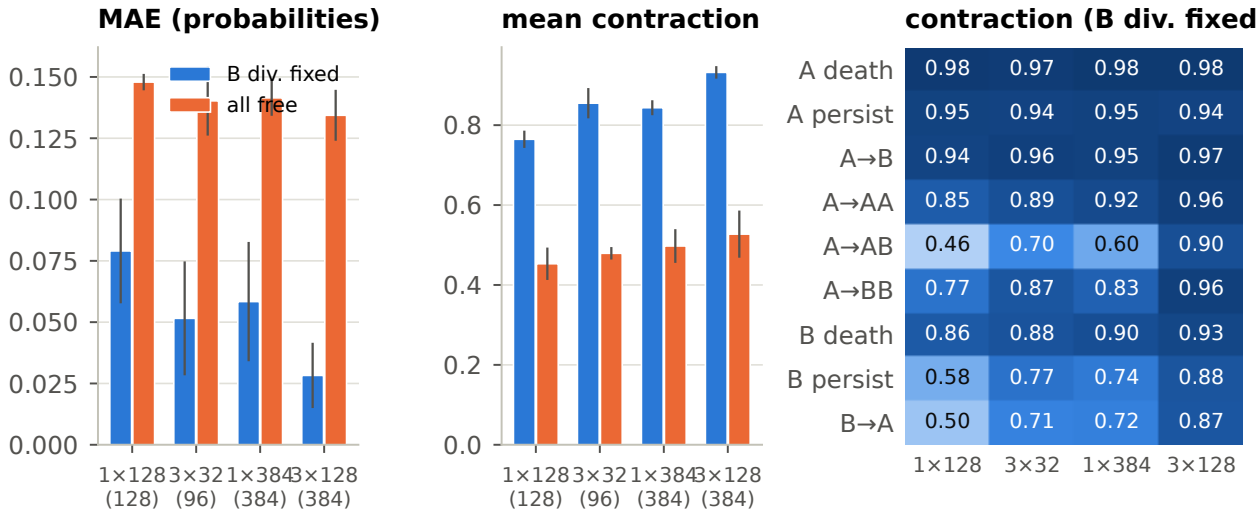


Figure 7: Left, centre: MAE and mean contraction per design (mean ± sd over 5 seeds); total organoids in brackets. Right: contraction per parameter, B divisions fixed.

Design (`reversible_ab`)	organoids	MAE	contraction	90% cov.	B→A	B→A contr.	PPC
1 × 128	128	0.079	0.77	0.87	0.61 ± 0.17	0.50	0.000
3 × 32	96	0.052	0.86	0.96	0.71 ± 0.13	0.71	0.008
1 × 384	384	0.058	0.84	0.93	0.67 ± 0.12	0.72	0.000
3 × 128	384	**0.028**	**0.93**	0.93	**0.78 ± 0.09**	**0.87**	0.031
`two_type`, any design	96–384	0.13–0.15	0.45–0.53	0.69–0.78	0.18–0.25	≤ -0.03	≤ 0.014

Comment

- At equal organoids, 3 endpoints halve the error of 1 (0.028 vs 0.058). Gains are in dynamics: A→AB contraction 0.90 vs 0.60, B→A 0.87 vs 0.72.
- 96 organoids over 3 endpoints ≈ 384 at one endpoint.
- With all transitions free no design recovers B's row (B→A contraction ≤ 0) — consistent with E0.
- Calibration fine everywhere (SBC 0.02–0.03). Caveat: GW, one truth, generations not days.

Leads to. One slice (GW, generations 4/8/12, 20k simulations). E4 maps the full grid of organoids × endpoint sets, including the real-like {6, 10}: GW first (§7), then BH (§8).

7 E4 — Amount of data: organoids and endpoints (GW)

Why this next. The design comparison (§6) showed that three endpoints beat one at equal organoids. For a plan we need the minimum number of organoids per endpoint for each parameter and the best endpoint set, systematically. GW comes first: it has no lifetimes, so it shows what endpoint counts can deliver when the time scale is known.

Question

How do contraction and error depend on organoids per endpoint (16–4096) and on the measured generations? What is the minimum number to recover each parameter when there is no lifetime ridge?

Setup

Simulator	Galton–Watson (endpoints in *generations*; day-mapping nominal)
Law (free)	`reversible_ab` (7 probabilities)
Data vs inference	well specified, $\theta^\star$
Design	clones per endpoint $\in \{16, 32, \ldots, 512\}$ and extensions $\{1024, 2048\}$, $\{4096\}$ × generations {10}, {6,10}, {6,10,14}, {4,6,10}
Inference / budget	NPE amortised, 5k; budget check: 20k at 4096 clones, {4,6,10}, same 5 datasets (§8.1)
Seeds / runs	5 / 180 = 120 (16–512) + 40 (1024, 2048) + 20 (4096); + 5 at 20k; all with 2 000 posterior samples, 200 SBC cases
Folder	`results/e4/gw/` (16–512), `results/e4/gw_large/` (1024, 2048; commit `5cb6018`), `results/e4/gw_xxl/` (4096; `07fb7fe`), `results/e4/gw_xxl_20k/` (20k; `07fb7fe`, preset docs `503ec72`); pooled tables `results/e4/gw_all_summary_*.csv`

Result

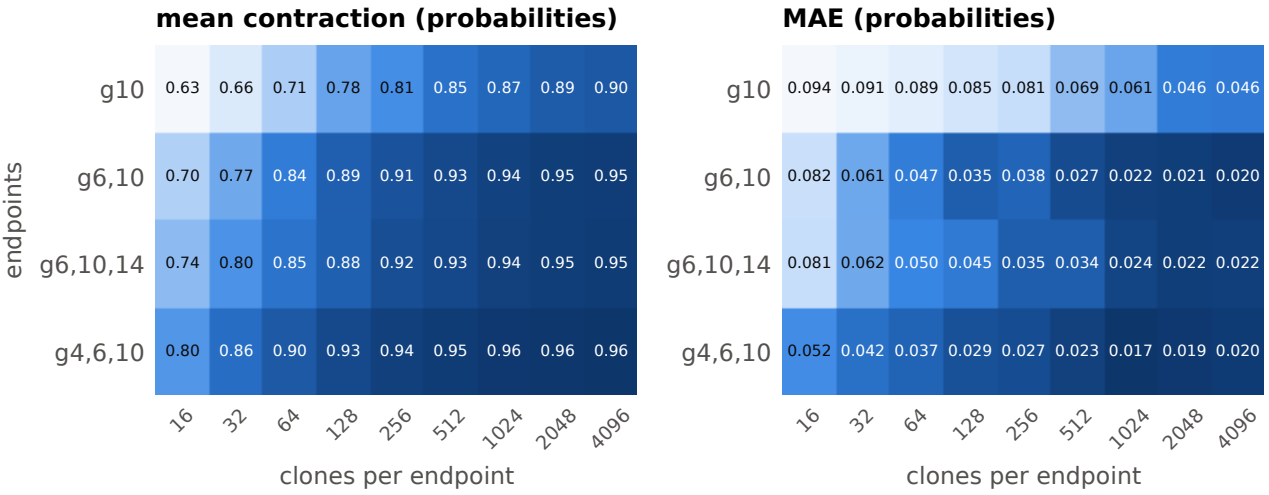


Figure 8: E4 GW: mean contraction and MAE of the outcome probabilities (mean over 5 seeds); endpoints in generations; NPE 5k; 1024–4096 from the extensions.

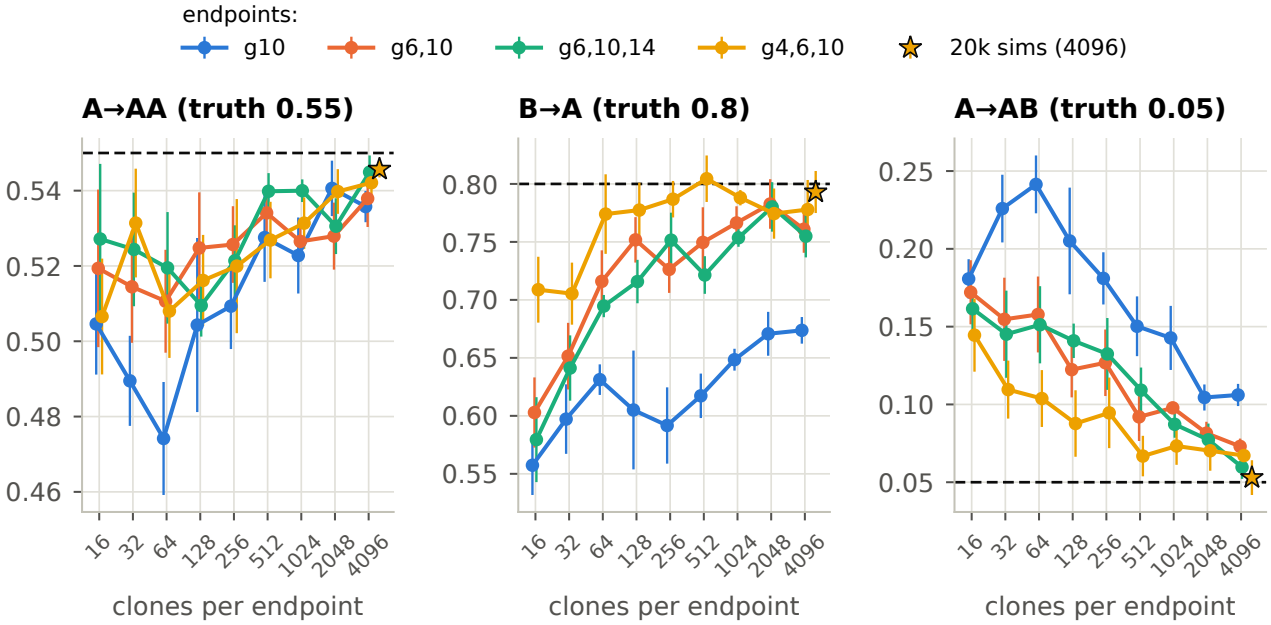


Figure 9: E4 GW: posterior means (mean ± SE over seeds) vs organoids per endpoint, NPE 5k; dashed = truth. Star: {4,6,10} at 4096 with 20k simulations (§8.1).

Minimum organoids per endpoint for "recovered" (contraction > 0.5, |error| < 0.05, ≥ 4/5 covered; must hold at every larger size up to 4096; — = not reached; NPE 5k; for BH see §8)

	GW, generations			
Parameter	{10}	{6, 10}	{6, 10, 14}	{4, 6, 10}
A death, A persist, A→B	16	16	16	16
B death	32	16	16	16
A→AA	256	16	16	2048[a]
A→BB	—	256	2048	**32**
B persist	—	512	—[c]	**64**
B→A	—	1024	—[c]	**128**
A→AB	—	512	2048	4096[b]

[a] error ≤ 0.04 at every size, but coverage 2–3/5 at 256–1024 (narrow posteriors, sd ≈ 0.02).
[b] also met at 128 and 512 (1024/2048: error 0.03, 3/5); at 4096 error 0.017, 4/5.
[c] 1024 up to 2048; at 4096 error 0.051 (B persist) and 0.053 (B→A), just above the 0.05 threshold.

Extensions: 1024, 2048, 4096 clones per endpoint (5 seeds, NPE 5k; MAE of probabilities / mean contraction / SBC coverage error)

Generations	512	1024	2048	4096
{10}	0.069 / 0.85 / 0.034	0.061 / 0.87 / 0.039	0.046 / 0.89 / 0.038	0.046 / 0.90 / 0.043
{6, 10}	0.027 / 0.93 / 0.037	0.022 / 0.94 / 0.044	0.021 / 0.95 / 0.045	0.020 / 0.95 / 0.050
{6, 10, 14}	0.034 / 0.93 / 0.041	0.024 / 0.94 / 0.047	0.022 / 0.95 / 0.055	0.022 / 0.95 / 0.058
{4, 6, 10}	**0.023** / **0.95** / 0.044	**0.017** / **0.96** / 0.052	**0.019** / **0.96** / 0.059	**0.020** / **0.96** / 0.064

Comment

- **Design ranking**: an early endpoint is worth most ({4,6,10} best at every size); a late extra endpoint ({6,10,14}) adds little over {6,10}; one endpoint plateaus and cannot separate A→AB from A→BB (0.15–0.24 vs 0.05; 0.15–0.20 vs 0.29 up to 512; still 0.11 and 0.24 at 4096).
- **Without lifetimes everything is recoverable**: with {4,6,10} and ≈128 organoids per endpoint all probabilities except A→AB (met at 128 and 512, not stably) meet the criterion; A→AA is right from 16 organoids. The GW numbers are the attainable target *if* division times are known (BH with lifetimes fixed: MAE 0.059, §9).
- **Plateau at ≈ 1024 per endpoint with NPE 5k** (extensions): multi-endpoint designs reach MAE ≈ 0.02 and contraction ≈ 0.95 and gain nothing from 2048 or 4096; with two or three endpoints every parameter meets the criterion at 1024–2048 except A→AB with {4,6,10} (intermittent). B→A contraction 0.87–0.93.
- **The plateau is the network's**: with 20k simulations at 4096 ({4,6,10}, same datasets) MAE 0.020 → 0.017, B→A error 0.047 → 0.032 (5/5 covered), SBC 0.064 → 0.044, support 0.59 → 0.78; training used 190 of 200 epochs on average (§8.1, after the BH grid).

- **A single endpoint never catches up**: it keeps improving, but at 4096 organoids (MAE 0.046) it is still worse than three endpoints at 32 (0.042); B→A, A→AB, A→BB and B persist never meet the criterion.
- **SBC error grows with size** (0.03 → 0.05–0.06 at 1024–4096) while coverage of the truth stays 0.87–0.98: at these sizes the amortised NPE (5k simulations) limits accuracy, as the 20k check confirms.
- Checks: single well-specified GW runs reach PPC 0.10 and support mass 0.07 ({4,6,10}, seed 4): borderline single values are not evidence of misspecification.
- Caveat: generations are not days; one truth, perfect assay, single A founders, 5k simulations (20k only for {4,6,10} at 4096). Planned extension (not run): `law=two_type` at 256–512 clones.

Leads to. The same grid in BH (§8), the simulator closer to the real data (time in hours, free lifetimes). Does the ranking survive, and are the GW numbers reached?

8 E4 — Amount of data: organoids and endpoints (BH), and comparison with GW

Why this next. GW (§7) recovers every probability by ≈128 organoids with {4,6,10}. BH adds two free lifetimes and measures time in days; it is the model that will be fitted to real data.

Question

As §7 for Bellman–Harris (endpoints in hours): how do contraction and error depend on organoids per endpoint (16–4096) and on the measured days? What is the minimum number to recover each parameter, and are the GW numbers reached?

Setup

Simulator	Bellman–Harris, Gamma lifetimes shape 8
Law (free)	`reversible_ab`: 7 probabilities + τ_A, τ_B (9)
Data vs inference	well specified, $\theta^\star$
Design	clones per endpoint ∈ {16, 32, 64, 128, 256, 512} and extensions {1024, 2048}, {4096} × endpoints {240}, {144, 240}, {144, 240, 336}, {96, 144, 240} h = days {10}, {6,10}, {6,10,14}, {4,6,10}
Inference / budget	NPE amortised, 5k; budget check: 20k at 4096 clones, days 4/6/10, same 5 datasets
Seeds / runs	5 / 180 = 120 (16–512) + 40 (1024, 2048) + 20 (4096); + 5 at 20k
Varied	clones × endpoint set; simulation budget (check only)
Folder	`results/e4/bh/` (16–512), `results/e4/bh_large/` (1024, 2048; commit `5cb6018`), `results/e4/bh_xxl/` (4096; `07fb7fe`), `results/e4/bh_xxl_20k/` (20k; `07fb7fe`, preset docs `503ec72`); pooled tables `results/e4/bh_all_summary_*.csv`

Result

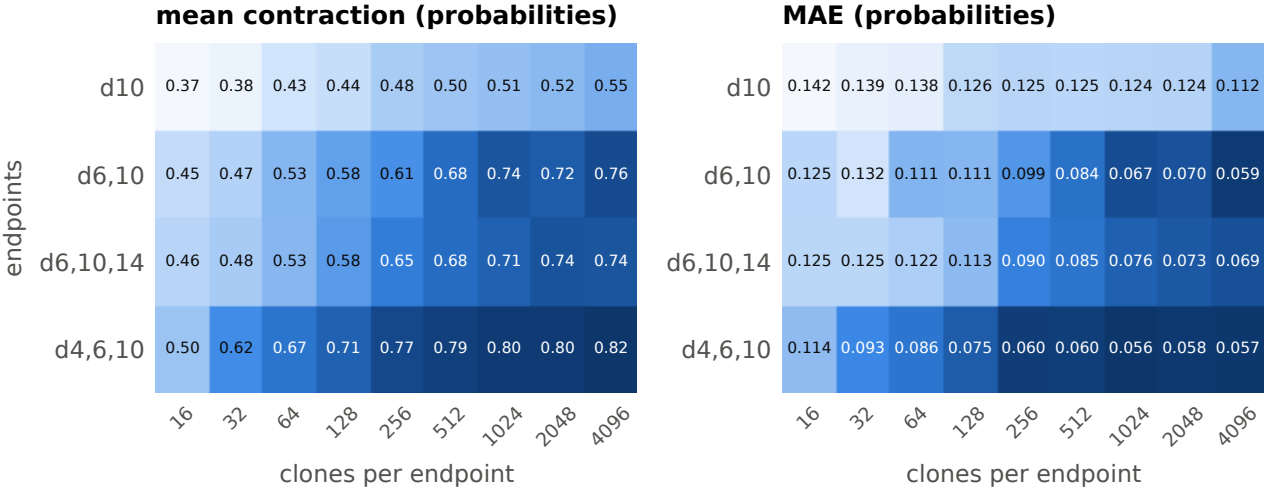


Figure 10: Mean contraction and MAE of the 7 outcome probabilities (mean over 5 seeds, NPE 5k); 1024–4096 from the extensions.

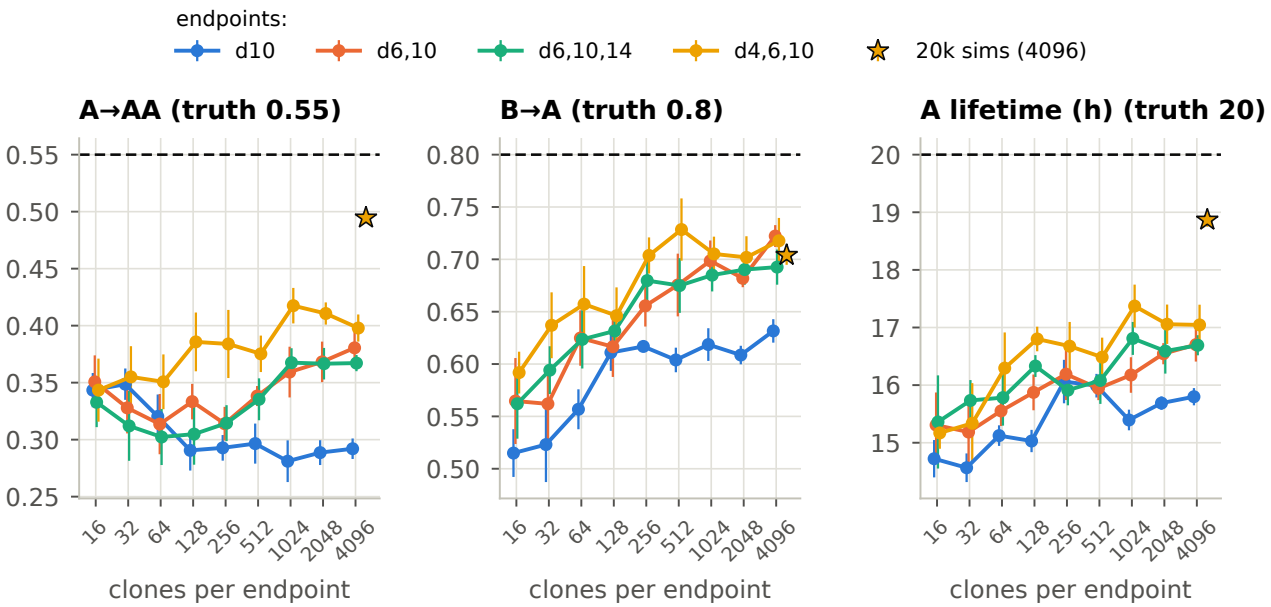


Figure 11: Posterior means (mean ± SE over seeds) vs organoids per endpoint, one line per endpoint set, NPE 5k; dashed = truth. Star: days 4/6/10 at 4096 with 20k simulations (§8.1).

Minimum organoids per endpoint for "recovered" (contraction > 0.5, |error| < 0.05, ≥ 4/5 covered); — = not reached at 4096 (NPE 5k; the 1024–4096 extensions change one entry).

Parameter	day 10	days 6, 10	days 6, 10, 14	days 4, 6, 10
A death, A→B (small truths)	16	16	16	16
A→BB	—	512	256	**128**
B death	—	256	256	**128**
A→AB	—	4096[a]	—	**256**
A→AA, A persist, B persist, B→A	—	—	—	—
τ_A, τ_B (biased: τ_A 15–17 h, 0/5 covered)	—	—	—	—

[a] met only at the largest size (error 0.026, 5/5); not a stable threshold.

Extensions: 1024, 2048, 4096 clones per endpoint (5 seeds, NPE 5k; probabilities: MAE, contraction; lifetimes: mean relative error)

Endpoints (days)	MAE 1024 / 2048 / 4096	contraction 1024 / 2048 / 4096	rel. error τ_A, τ_B 1024 / 2048 / 4096	A→AA (0.55) 512 → 2048 → 4096	τ_A (20 h) 2048 / 4096
{10}	0.124 / 0.124 / 0.112	0.51 / 0.52 / 0.55	0.32 / 0.30 / 0.30	0.30 → 0.29 → 0.29	15.7 / 15.8
{6, 10}	0.067 / 0.070 / 0.059	0.74 / 0.72 / 0.76	0.29 / 0.27 / 0.24	0.34 → 0.37 → 0.38	16.5 / 16.7
{6, 10, 14}	0.076 / 0.073 / 0.069	0.71 / 0.74 / 0.74	0.25 / 0.27 / 0.27	0.34 → 0.37 → 0.37	16.6 / 16.7
{4, 6, 10}	**0.056** / **0.058** / **0.057**	**0.80** / **0.80** / **0.82**	0.24 / 0.27 / 0.27	0.38 → 0.41 → 0.40	17.1 / 17.0

A→AA and τ_A covered 0/5 at every size and design with NPE 5k.

Comparison with GW (§7; same grid, endpoints in generations)

	BH (days)	**GW** (generations)
Contraction, {4,6,10}, 16/512/2048/4096	0.50 / 0.79 / 0.80 / 0.82	0.80 / 0.95 / 0.96 / 0.96
MAE, {4,6,10}, same sizes	0.114 / 0.060 / 0.058 / 0.057	0.052 / 0.023 / 0.019 / 0.020
One late endpoint, contraction	0.37 / 0.50 / 0.52 / 0.55	0.63 / 0.85 / 0.89 / 0.90 (A→AB/BB conf.)
A→AA (0.55), NPE 5k	0.28–0.42, 0–2/5 covered	0.47–0.55 at every design
A→AA, 4096, {4,6,10}, 5k / 20k	0.40 (0/5) / 0.50 (3/5)	0.54 / 0.55 (5/5)
B→A (0.80), {4,6,10}, 128/512/2048/4096	0.65 / 0.73 / 0.70 / 0.72	0.78 / 0.80 / 0.77 / 0.78
SBC error, 16–512 / 1024–4096	0.025–0.034 / 0.029–0.039	0.024–0.044 / 0.038–0.064
Support mass / PPC (per run)	0.37–0.92 / ≤ 0.042	0.07–0.89 / ≤ 0.10 (groups ≤ 0.021)

Comment

- **Same design ranking as GW** (§7): the ranking is a property of the endpoint information, not of the time model. An early endpoint (day 4) helps most: days 4/6/10 are best on every metric at every size and the only design where A→AB is recovered from a moderate size (0.054 at 2048, 0.042 at 4096, truth 0.05; days 6/10 meet it only at 4096). Day 10 alone plateaus (contraction 0.37 → 0.50 → 0.52 → 0.55 from 16 to 512 to 2048 to 4096).

- With several endpoints and NPE 5k more organoids help up to ≈512–1024 (MAE 0.114 → 0.060 → 0.056 for days 4/6/10; 0.058 at 2048, 0.057 at 4096). B→A rises 0.59 → 0.73 and is covered at every size, but not recovered (0.61–0.72 at 1024–4096).
- **With NPE 5k the ridge does not yield to sample size.** At 4096 clones per endpoint (up to 12 288 organoids) A→AA is 0.29–0.40 (truth 0.55) and τ_A 15.8–17.0 h (truth 20), 0/5 covered in every design. The posterior creeps along the ridge (days 4/6/10: 0.34 → 0.41 from 16 to 2048, flat at 4096) with τ_A contraction 0.99: confidently on the wrong point.
- **But this is a training limit, not a data limit** (§8.1): with 20k simulations on the same 4096-clone datasets A→AA moves 0.40 → 0.50 and τ_A 17.0 → 18.9 h (3/5 covered each). At realistic sizes (≤ 512, untested at 20k) lifetime information remains the practical lever (§9, X5b).
- Calibration at 1024–4096: SBC error rises mildly (0.028–0.039), PPC ≤ 1.4%, support mass 0.37–0.81.
- "Recovered" for A death / A→B is easy because the truths are small (|error| < 0.05 is loose there).
- GW, without lifetimes, reaches contraction 0.80 → 0.95 → 0.96 (16 → 512 → 2048, also 0.96 at 4096) and recovers A→AA at every design (table above).
- Caveat: one truth, perfect assay, single A founders, 5k simulations (20k only for one design at 4096).

8.1 Simulation budget check: 5k vs 20k simulations at 4096 clones (BH and GW)

Question

At the largest designs, is the GW plateau (§7) and the stuck BH ridge (above) a limit of the data or of the 5k-simulation network?

Setup

Design	4096 clones per endpoint; best design per simulator: GW generations {4,6,10}, BH days {4,6,10} = {96,144,240} h
Inference / budget	NPE amortised, 5k (from `*_xxl`) vs 20k (`budget.simulations=20000`); epoch cap 200
Seeds / runs	5 paired seeds (same datasets) / 5 per simulator at 20k
Folder	`results/e4/gw_xxl_20k/`, `results/e4/bh_xxl_20k/` (commit `07fb7fe`; preset docs `503ec72`); table `report/tables/e4_budget.csv`

Result

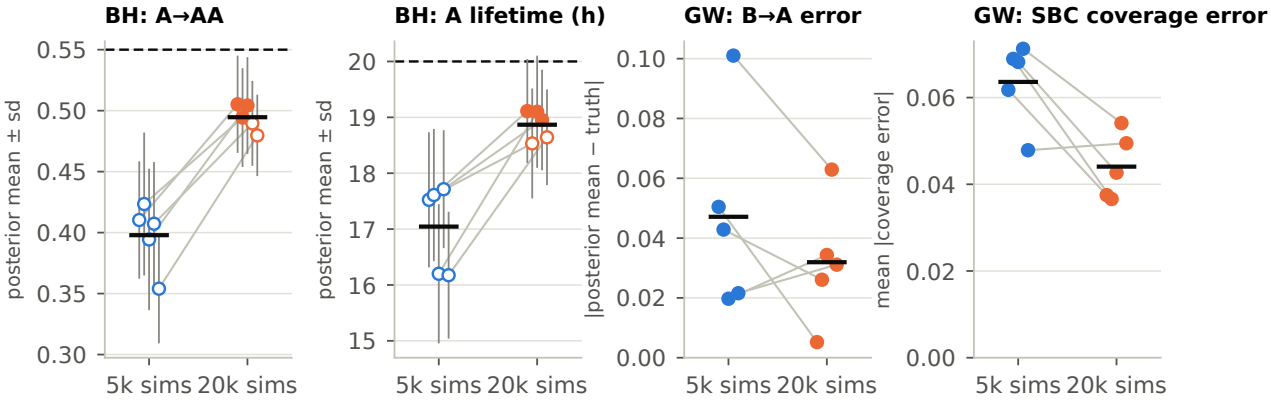


Figure 12: 5k vs 20k training simulations on the same 5 datasets (lines join seeds). BH, days 4/6/10: A→AA and τ_A posterior mean ± sd; filled = truth in the 90% interval, open = not covered; dashed = truth. GW, generations 4/6/10: B→A absolute error and run SBC coverage error. Bars = mean over seeds.

	GW {4,6,10}		BH {96,144,240} h	
	5k	20k	5k	20k
MAE (probabilities)	0.020	**0.017**	0.057	**0.049**
Rel. error lifetimes	–	–	0.27	**0.15**
Mean contraction (all parameters)	0.964	0.977	0.842	0.893
Coverage (run, 90%)	0.93	0.93	0.76	**0.93**
SBC coverage error	0.064	**0.044**	0.039	0.040
Support mass	0.59	0.78	0.56	0.51
Epochs, mean (cap 200)	162	190	160	**201** (all 5)
B→A (0.80): mean, error, covered	0.78, 0.047, 4/5	0.79, 0.032, 5/5	0.72, 0.082, 5/5	0.70, 0.096, 5/5
A→AA (0.55): mean, covered	0.54, 5/5	0.55, 5/5	0.40, 0/5	**0.50, 3/5**
τ_A (20 h): mean, covered	–	–	17.0, 0/5	**18.9, 3/5**
τ_B (30 h): mean, covered	–	–	18.4, 2/5	22.5, 5/5

Comment

- **The 5k network, not the data, sets the large-design plateau.** Four times the training simulations lower the error on the same datasets in both simulators and improve calibration (GW SBC 0.064 → 0.044) or coverage (BH 0.76 → 0.93). At 5k, 2048 → 4096 is flat for both simulators.
- **The BH ridge moves with enough data when the network is trained enough**: A→AA 0.40 → 0.50 and τ_A 17.0 → 18.9 h, each 3/5 covered (0/5 at 5k); lifetime error 0.27 → 0.15. "The ridge does not yield to sample size" holds for NPE at 5k simulations, not for the data.
- Training stopped at the 200-epoch cap in all 5 BH runs (and 190 on average in GW), so 20k is probably not the limit either: 5k (and possibly 20k) results at large sizes are lower bounds on attainable accuracy.
- Not tested: 20k at realistic sizes ($\leq$ 512 per endpoint). Whether the ridge moves there is the decisive follow-up (§28); until then lifetime information is the practical fix at realistic sizes.
- B→A in BH does not improve (0.72 → 0.70, error 0.082 → 0.096, still covered); BH support mass slightly lower (0.56 → 0.51).
- Caveat: one design per simulator, 5 seeds, one truth.

Leads to. With the standard 5k network A→AA and τ_A are wrong at *every* size (16–4096) and endpoint set, with coverage falling as data grow; only a 20k network at 4096 clones moves them most of the way. At realistic sizes this is a bias that more organoids do not remove. Part III opens with the hypothesis that explains it (§9).

Part III

Bellman–Harris timing: the lifetime ridge

Motivation. E4 BH (§8) left a bias that more organoids do not remove with the standard 5k-simulation network, not even 4096 per endpoint (a 20k network at 4096 moves it most of the way, §8.1; at realistic sizes that is untested). This part identifies it as a ridge (§9), asks whether it is a property of the hand-picked truth (X5d), and tests the remedy suggested by the ridge, lifetime information, in its soft (literature prior, X5b) and hard (fixed lifetimes, E3a) form. Fixing lifetimes rests on timing assumptions: one value per type (E3a), the lifetime variability (E3b), and the absence of quiescence (E3c), which are tested last.

9 BH ridge — lifetime vs self-renewal

Why this next. E4 BH (§8): A→AA 0.28–0.42 and τ_A 15–17 h at every design up to 4096 organoids per endpoint with NPE 5k, 0/5 covered, whereas GW recovers A→AA at every design (§7). The only structural difference is the free time scale. Hypothesis: endpoint counts fix the growth rate, which faster cycles or more self-renewal produce equally.

Question

Is the BH baseline bias (A→AA ≈ 0.32, $\tau_A \approx 16$ h) a ridge on which endpoint counts fix only a combination of lifetime and self-renewal, with the prior picking the point?

Setup

Simulator	Bellman–Harris, shape 8
Law (free)	control: `reversible_ab` (9); lifetimes fixed: 7; flat prior: 9
Data vs inference	data always $\theta^\star$; only the inference model changes
Design	real-like: 144/240 h × 165/101
Inference / budget	NPE amortised, 5k
Arms	control (from E3b, shape 8); τ_A=20, τ_B=30 h fixed; flat prior on A→AA (concentration 1 vs 0.2 for A's other outcomes: A→AA ∼ Beta(1,1))
Seeds / runs	5 / 10 (+5 control)
Folder	`results/bh_ridge/` (control: `results/e3b/*lifetime_shape=8*`)

Result

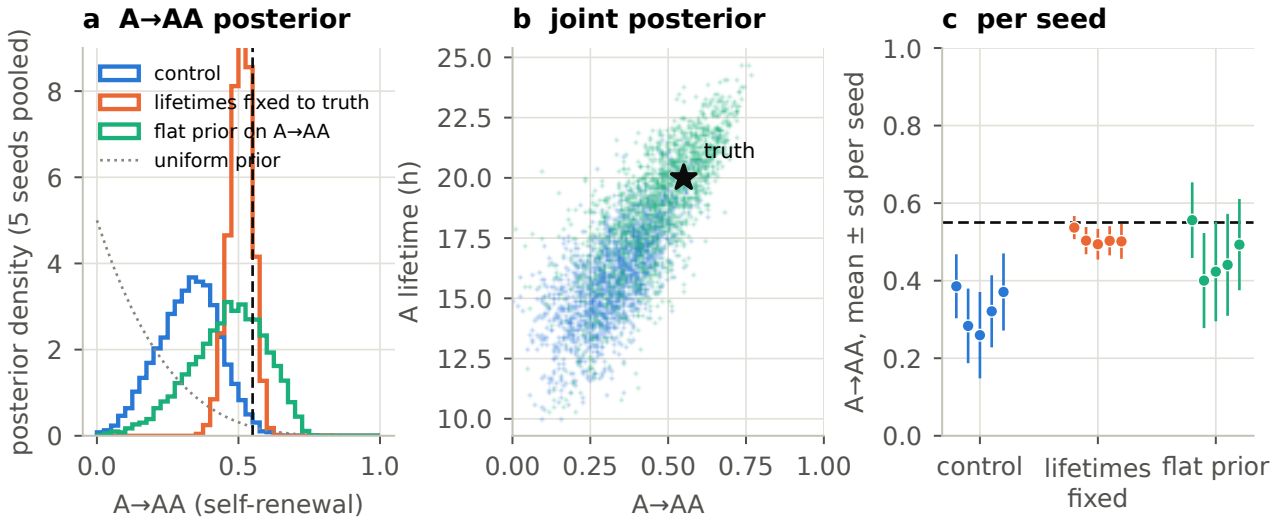


Figure 13: (a) A→AA posterior pooled over 5 seeds; dotted: uniform-Dirichlet prior Beta(1,5). (b) Joint posterior of A→AA and τ_A (blue: control; green: flat prior); star = truth. (c) Per-seed means ± sd; dashed = truth.

	Control	Lifetimes fixed	Flat prior
A→AA (0.55)	0.32 ± 0.10, 0/5, contr. 0.52	**0.51 ± 0.04, 5/5, 0.93**	0.46 ± 0.12, 5/5, 0.83
A→AB (0.05)	0.23, 3/5	0.17, 2/5	0.20, 5/5, contr. 0.00
B→A (0.80)	0.63, 5/5	0.68, 5/5	0.46, 3/5
τ_A (20 h)	15.9 ± 1.9, 0/5	fixed	18.9 ± 2.2, 5/5
MAE (probabilities)	0.109	**0.059**	0.117
SBC error / PPC / support	0.025 / 0.008 / n/r	0.031 / 0.025 / 0.65	0.079 / 0.000 / 0.51

Comment

- Ridge confirmed: posterior corr(A→AA, τ_A) ≈ +0.7 (E3c). Removing the lifetime degree of freedom recovers self-renewal and halves the error.
- The prior chose the point: only 1.8% of uniform-Dirichlet mass lies above 0.55. A flat prior re-centres (5/5 covered) but stays wide, degrades B→A and calibration: not a fix.
- GW has no lifetimes and no ridge: the GW control at the same design gives A→AA 0.52 (§15), and 0.47–0.55 at every E4 design (§7).
- Sample size alone, with the standard 5k-simulation NPE, does not leave the ridge: A→AA stays 0.29–0.41 and τ_A 15.7–17.1 h (0/5) up to 4096 clones per endpoint in every E4 design. With 20k simulations at 4096 clones (days 4/6/10) it does move most of the way: A→AA 0.50, τ_A 18.9 h (3/5 each; §8.1). The data can resolve the ridge at very large designs if the network is trained enough; whether they do at realistic sizes ($\leq$ 512) with 20k is untested.
- Consequence: at realistic sizes measure A's division time (live imaging) or use it as an informative prior; otherwise report BH results with prior sensitivity, and train BH fits on $\geq$ 20k simulations. X5 tested both: a right ±25% lifetime prior gives A→AA 0.46, MAE 0.068 (§11); averaged over prior-drawn truths A→AA is learned well (contraction 0.82, §10), so the ridge bites at tail truths such as $\theta^\star$.

Leads to. Two questions: is the ridge bias a property of this truth, whose A→AA has only 1.8% prior mass above it (X5d, §10)? And can partial lifetime information — a literature prior — replace exact lifetimes (X5b, §11)?

10 X5d — Role of the prior III: performance averaged over the prior (GW and BH)

Why this next. The three failures so far — E0's B row (§3), the BH ridge (§9), and later the original E2c (§24) — all occurred at truths in the prior tail. X5a (§4) showed that the B-row failure is structural. Is the BH ridge equally structural, or does BH work for typical truths?

Question

Every other experiment uses one hand-picked truth, partly in the prior tail. Averaged over truths drawn from the prior, are GW and BH calibrated and informative?

Setup

Data = inference	`reversible_ab`, 30 truths drawn from the prior (`truth.source=prior`, `data_seed` = seed 0–29; no redraws needed); BH lifetime prior log-uniform [10, 80] h
Design	real-like: GW generations 6/10, BH 144/240 h; 165/101 clones
Seeds / runs	30 per simulator (60)
Prior-mean baseline	error of the uniform Dirichlet mean (1/6 for A's outcomes, 1/3 for B's)
Folder	`results/x5/d_gw/`, `results/x5/d_bh/` (`prior_average_parameters.csv`, `prior_average_runs.csv`)

Result

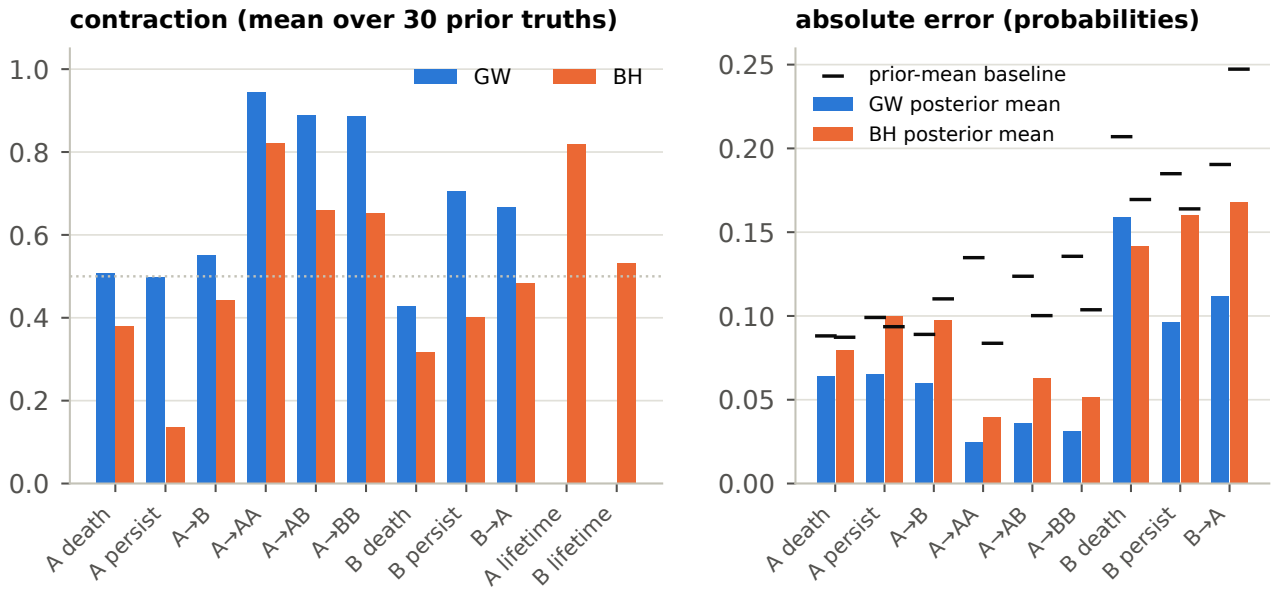


Figure 14: X5d: per-parameter contraction (left) and absolute error of the posterior mean (right; black ticks = prior-mean baseline), averaged over 30 truths drawn from the prior; GW vs BH.

	GW	BH
SBC coverage error	0.032	0.028
90% coverage across truths	0.90	0.90
Mean contraction	0.68	0.51
MAE probabilities (prior-mean baseline)	0.072 (0.139)	0.100 (0.129)
A→AA contraction	0.94	0.82
B→A contraction / coverage	0.67 / 0.90	0.48 / 0.80
Support mass median (min)	0.88 (0.54)	0.89 (0.29)
PPC outside 95%	0.006	0.003

BH mean contraction includes τ_A (0.82) and τ_B (0.53); relative lifetime error 0.21 / 0.35. Numbers: `tables/x5d.csv`.

Comment

- **The methods are calibrated on average**: with truths drawn from the prior both simulators give 90% coverage and SBC error $\approx$ 0.03. Single-truth failures are about specific truths and priors, not broken inference.
- GW learns A's divisions well (contraction 0.89–0.94), A's death/persist/switch about halfway ($\approx$ 0.5), B's row moderately (0.43–0.71; B death weakest).
- **BH learns less** (contraction 0.51 vs 0.68; MAE 0.100 vs 0.072): the lifetimes absorb information; A persist barely moves (contraction 0.14, error no better than the prior mean) and B death / B persist stay close to the prior-mean baseline (0.142 / 0.160 vs 0.170 / 0.164).
- **The BH ridge is a tail-truth effect**: averaged over the prior A→AA is learned well (contraction 0.82), so the failure at $\theta^\star$ (A→AA 0.55, 0/5 covered, §9) is specific to a truth with 1.8% prior mass above it, where the prior picks the point on the lifetime/self-renewal ridge (at realistic sizes; at 4096 clones per endpoint and 20k simulations the data largely override it, §8.1).
- Caveat: X5d uses real-like generations 6/10; the E0/design-comparison single-truth runs used 4/8/12, so they are not directly comparable.

Leads to. The ridge bites where the prior chooses the point. The natural fix is to give the prior information about lifetimes (X5b, §11).

11 X5b — Role of the prior IV: informative lifetime prior, right vs wrong (BH)

Why this next. Fixing lifetimes to the truth removes the ridge (§9), and X5d (§10) showed that the bias arises where the vague lifetime prior picks the point on it. Literature division times are available, but only approximately, and may be wrong.

Question

Fixing lifetimes resolves the BH ridge (§9). What does an informative (literature) lifetime prior do if it is centred on the truth, or 50% too long (truth outside the prior)?

Setup

Simulator / law	Bellman–Harris shape 8, `reversible_ab` (9 free); data $\theta^\star$ (τ_A 20 h, τ_B 30 h)
Lifetime priors (log-uniform)	vague: [10, 80] h for both (= BH ridge control, bit-identical to E3b); right (±25%): A [16, 25], B [24, 37.5] h; wrong (+50%): A [24, 37.5], B [36, 56.25] h
Seeds / runs	5 / 15
Folder	`results/x5/b_vague/`, `b_right/`, `b_wrong/` (`x5b_groups.csv`)

Result

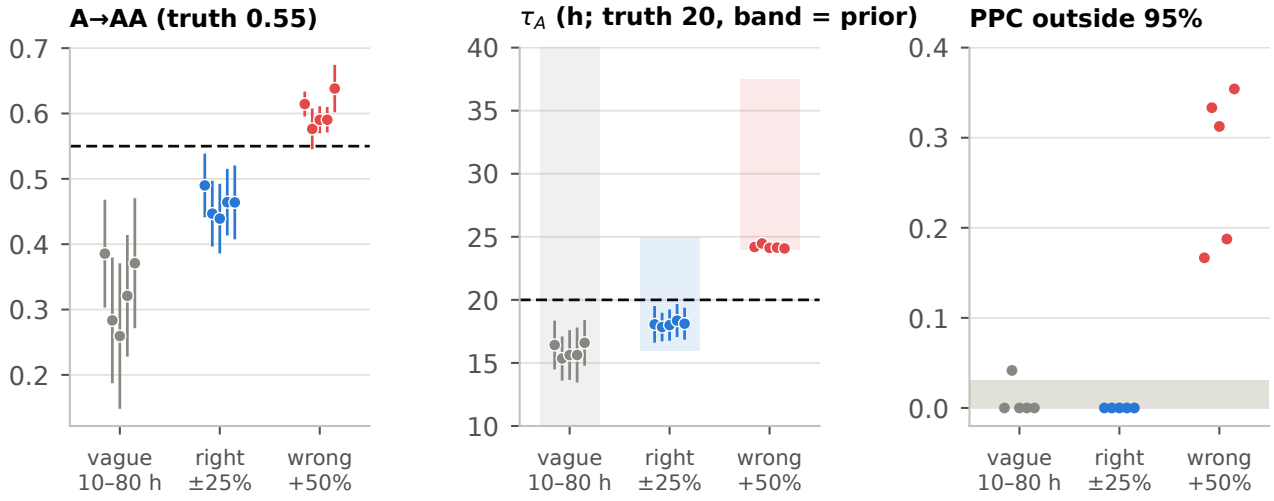


Figure 15: X5b: A→AA and τ_A per seed (posterior mean ± sd) and PPC per run for the three lifetime priors. Dashed = truth; shaded band = prior range of τ_A (vague prior extends to 80 h); grey band in the PPC panel = well-specified range.

Parameter (truth)	vague [10, 80] h	right (±25%)	wrong (+50%)
A→AA (0.55)	0.32 ± 0.10, 0/5	0.46 ± 0.05, 2/5	0.60 ± 0.03, 1/5
A→AB (0.05)	0.23 ± 0.12, 3/5	0.18 ± 0.08, 2/5	0.20 ± 0.06, 1/5
A→BB (0.29)	0.19 ± 0.09, 5/5	0.19 ± 0.06, 1/5	0.17 ± 0.04, 1/5
B→A (0.80)	0.63 ± 0.18, 5/5	0.69 ± 0.16, 5/5	0.68 ± 0.14, 5/5
τ_A (20 h)	15.9 ± 1.9, 0/5	18.1 ± 1.3, 4/5	24.2 ± 0.2 (prior edge), 0/5
τ_B (30 h)	20.2 ± 9.5, 4/5	29.5 ± 3.6, 5/5	49 ± 5.3, 0/5
MAE probabilities / rel. error lifetimes	0.109 / 0.26	**0.068** / 0.06	0.082 / 0.42
Support mass / PPC outside 95%	0.74–0.83 / 0.008	0.57–0.72 / 0.000	**0.00–0.02 / 0.27**
SBC coverage error	0.027	0.028	0.027

Mean over 5 seeds ± mean posterior sd, seeds covered. Lifetimes fixed exactly (§9): A→AA 0.51, MAE 0.059.

Comment

- **A right literature prior moves along the ridge towards the truth**: A→AA 0.32 → 0.46, lifetimes about right, MAE 0.109 → 0.068 — most of the way to fixing lifetimes exactly (0.51, 0.059), at the real-like design. More organoids alone reach a similar point only at 4096 per endpoint and with 20k simulations (A→AA 0.50, §8.1); at realistic sizes lifetime information is the practical lever.
- **A prior 50% too long is clearly flagged**: support ≈ 0, PPC 0.27 (0.17–0.35), A's lifetime pinned at the prior's lower edge (the data push against it). Consistent with E3a, where too-long lifetimes fail visibly and too-short ones bias silently. Probabilities remain off (1/5 covered) despite high contraction.
- Caveat: one truth ($\theta^\star$), 5 seeds; a too-short wrong prior was not tested.

Leads to. Lifetime information works. In practice it will be a measured, possibly pooled value fixed in the model: E3a (§12) tests per-type vs global and wrong fixed values.

12 E3a — Per-type division time vs one global value (BH)

Why this next. The ridge (§9) and X5b (§11) make A's division time the key measurement. The real division time is estimated globally, but B may divide more slowly than A. What happens if one pooled or slightly wrong value is fixed?

Question

If B divides more slowly than A, can both lifetimes be inferred? What happens if both are fixed to one global (pooled) value, or to a wrong one (±25%)?

Setup

Simulator	Bellman–Harris, shape 8
Law (free)	free arm: `reversible_ab` (9); fixed arm: lifetimes fixed to T (7)
Data model	$\tau_A = 20$ h, $\tau_B \in \{20, 30, 45\}$ h; outcome probabilities $\theta^\star$
Inference model	(i) both lifetimes inferred; (ii) both fixed to pooled $T = (20 + \tau_B)/2$; (iii) pooled ±25%
Design	real-like 144/240 h × 165/101
Inference / budget	NPE amortised, 5k
Seeds / runs	5 / 60 (3 τ_B × (1 free + 3 fixed))
Control	free arm at the same τ_B
Folder	`results/e3a/free`, `results/e3a/fixed`

Result

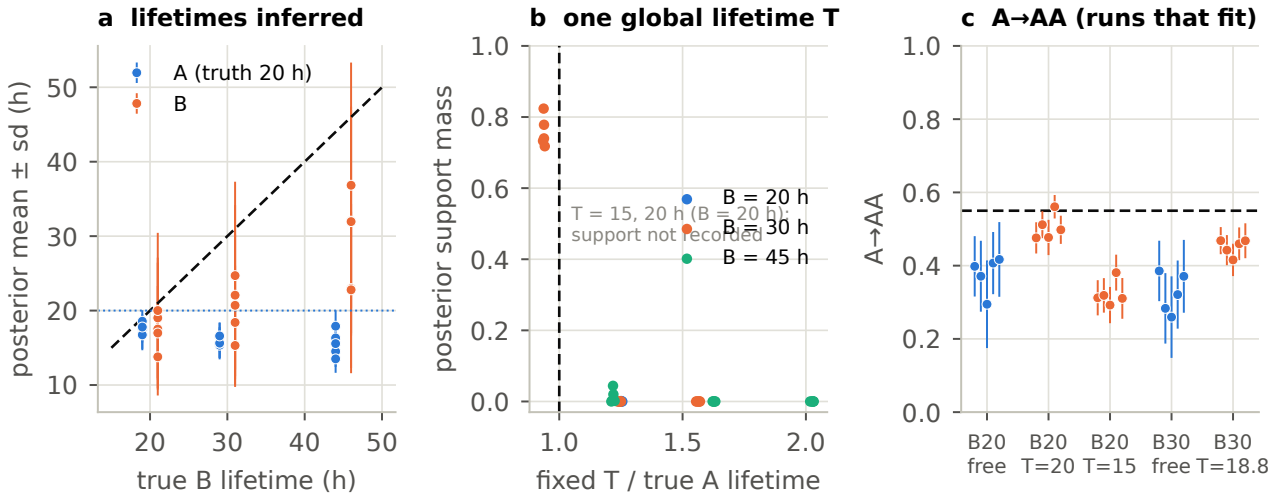


Figure 16: (a) Inferred lifetimes vs true τ_B (per run, ± sd); dotted = true τ_A. (b) Posterior support mass vs fixed T relative to the true τ_A (20 h). (c) A→AA in the settings that fit.

τ_B / T		A persist	A→AA	A→AB	B→A	MAE	PPC	support
20 / free	control	0.11	0.38 [1/5]	0.21	0.68	0.087	0.004	n/r
20 / **20**	T = truth	0.05	**0.50 [4/5]**	0.14	0.75	**0.042**	0.004	n/r
20 / 15	−25%	0.20 (+0.08)	0.32	0.15	0.79 (+0.10)	0.081	0.013	n/r
30 / 18.75	−25%	0.05	0.45 (+0.13)	0.13	0.59 [1/5]	0.072	0.000	0.76
any / ≥ 24.4	$T > \tau_A$	no estimate (0–2 000 samples; boundary values)				—	0.35–0.51	≈ 0

Free arm: τ_A 17.5 / 15.9 / 15.6 h and τ_B 17.5 / 20.2 / 29.3 h for true τ_B 20 / 30 / 45 (contraction of τ_B 0.84 / 0.75 / 0.49).

Comment

- Lifetimes fixed at the truth give the best BH recovery of all (A→AA 0.50, MAE 0.042, contraction 0.92) — direct evidence for the ridge.
- Inferred τ_A is biased short whatever τ_B; τ_B is compressed towards τ_A and weakly informed.
- T longer than A's lifetime ⇒ the model cannot fit (support ≈ 0, PPC 0.35–0.51): *detectable*. Depends on T vs τ_A, not on the pooled mean.
- T a little shorter than τ_A ⇒ fits normally but shifts outcomes by up to 0.10 (A persist absorbs the faster clock): *undetectable*.

- Consequence: measure A's division time per type; never use a pooled mean when B is slower. Caveat: coarse T grid (failure threshold between 20 and 24.4 h).

Leads to. Fixing the mean lifetime also fixes an assumed variability (Gamma shape 8, CV 0.35, in every BH run). E3b (§13) tests that assumption.

13 E3b — Desynchronisation: lifetime CV misspecified (BH)

Why this next. Every BH fit so far assumed lifetime shape 8 (CV 0.35); E3a (§12) showed that a wrong lifetime *mean* can bias silently. Real cells desynchronise; is a wrong *variability* equally harmful?

Question

How biased is the fit when the data's lifetime variability (Gamma shape 1–32, CV 1–0.18, or per type) differs from the assumed shape 8 (CV 0.35)? Does the PPC see it?

Setup

Simulator	Bellman–Harris
Law (free)	`reversible_ab` (9)
Data model	shape $\in \{1, 2, 4, 8, 32\}$ and per type {A 4, B 16}, {A 16, B 4}
Inference model	shape 8 always (one network per seed: paired shifts)
Design	real-like 144/240 h × 165/101
Inference / budget	NPE amortised, 5k
Seeds / runs	5 / 35
Control	data shape 8 (= BH baseline)
Folder	`results/e3b/`

Result

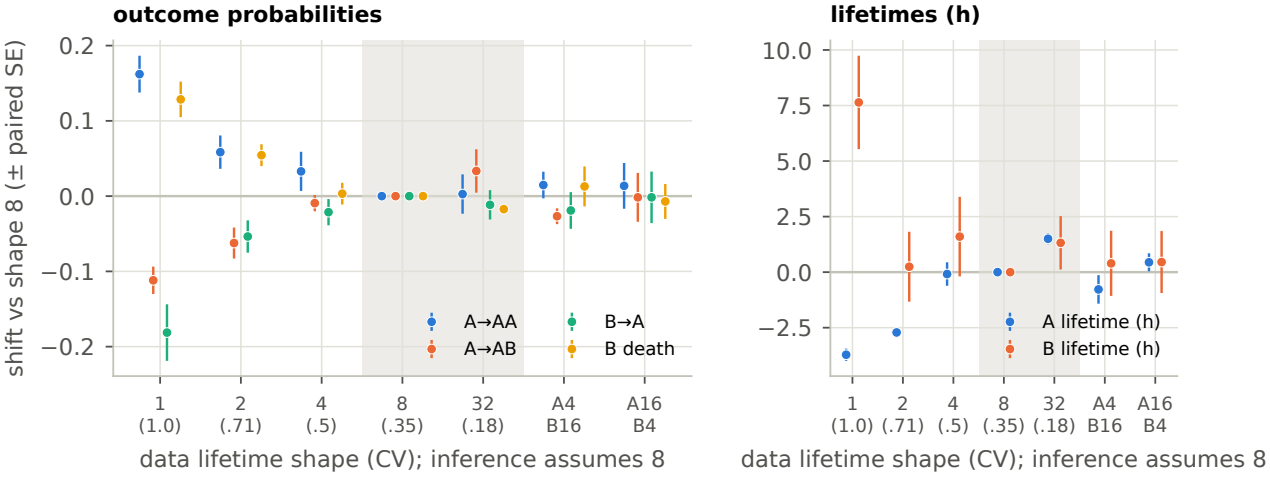


Figure 17: Shift vs shape 8 (± paired SE); grey band = CV 0.18–0.35, around the assumed value.

Data shape (CV)	A→AA	A→AB	B death	B→A	τ_A (h)	τ_B (h)	PPC
8 (0.35) control	0.32	0.23	0.17	0.63	15.9	20.2	0.008
1 (1.0)	+0.16 ± .02	−0.11	+0.13	−0.18± .04	−3.7± 0.3	+7.6 ± 2.1	0.004
2 (0.71)	+0.06	−0.06	+0.05	−0.05	−2.7	+0.2	0.000
4 (0.5)	+0.03	−0.01	+0.00	−0.02	−0.1	+1.6	0.000
32 (0.18)	+0.00	+0.03	−0.02	−0.01	+1.5 ± 0.2	+1.3	0.000
A 4 / B 16; A 16 / B 4	$\leq$ 0.01	$\leq$ 0.03	$\leq$ 0.01	$\leq$ 0.02	−0.8 / +0.4	$\leq$ 0.5	$\leq$ 0.004

Comment

- CV 0.18–0.5 (incl. per-type mixes) is harmless: shifts $\leq$ 0.035 and $\leq$ 1.6 h, small vs posterior sd ($\approx$ 0.1, 2 h).
- Near-exponential cycle times are read as a faster clock with a different division split; bias is large and invisible to the PPC. (A→AA moves *towards* the truth only because the control is biased low.)
- Consequence: a rough CV estimate (live imaging) suffices; if CV $\gtrsim$ 0.7 change the shape in the model. Shape was never inferred; support mass not recorded.

Leads to. One more parameter can slow growth without changing the clock: quiescence (A persist). E3c (§14) asks whether it lies on the same ridge.

14 E3c — Quiescence vs division time (BH)

Why this next. The ridge (§9) trades self-renewal against lifetime. Quiescent cells (A persist) also slow growth, so they may be a third direction on it; if so, fixing quiescence wrongly (e.g. to 0) would shift the lifetimes.

Question

Are quiescence (A persist) and A's lifetime confounded? What happens if quiescence is fixed to 0 while the data have A persist $\in \{0.03, 0.15, 0.30\}$?

Setup

Simulator	Bellman–Harris, shape 8
Law (free)	free arm `reversible_ab` (9); fixed arm A persist = 0 (8)
Data model	A persist 0.03 / 0.15 / 0.30, other A entries rescaled (A→AA 0.55 / 0.48 / 0.40)
Design	real-like 144/240 h × 165/101
Inference / budget	NPE amortised, 5k
Seeds / runs	5 / 30
Control	free arm at the same data
Folder	`results/e3c/persist_free`, `results/e3c/persist_fixed0`

Result

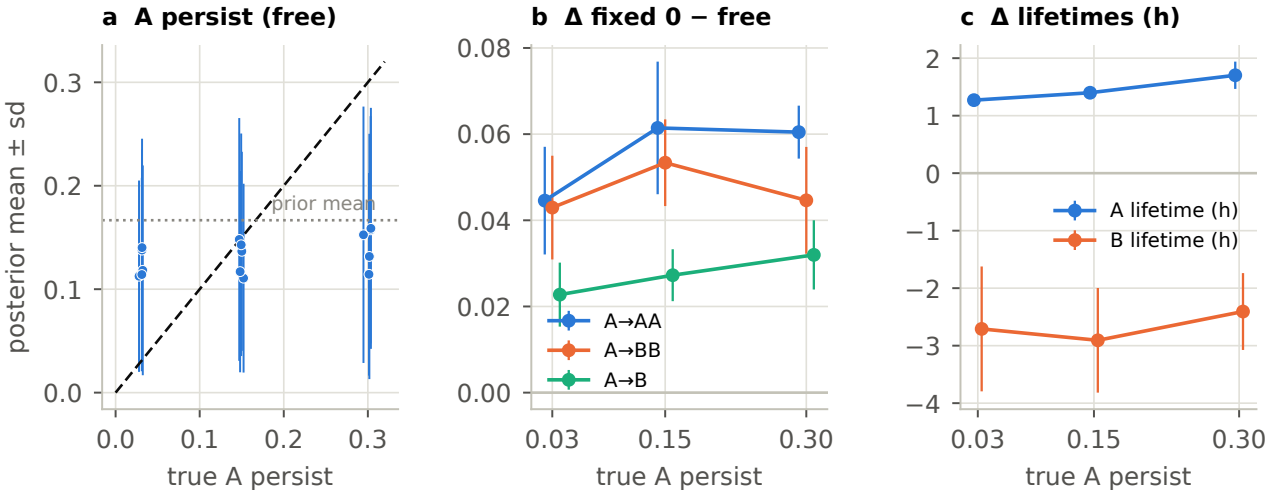


Figure 18: (a) A persist posterior (free) vs truth; dotted = prior mean 1/6. (b, c) Shift fixed-0 minus free (± paired SE).

True A persist	free: A persist (contr.)	free: τ_A	fixed 0: $\Delta\tau_A$	ΔA→AA	$\Delta\tau_B$	PPC
0.03	0.12 (0.51)	15.9	+1.3 ± 0.1	+0.05	−2.7	0.000
0.15	0.13 (0.47)	17.9	+1.4	+0.06	−2.9	0.000
0.30	0.14 (0.35)	19.3	+1.7	+0.06	−2.4	0.004

Posterior correlations (free arm): corr(A persist, τ_A) ≈ −0.6; corr(A→AA, τ_A) ≈ +0.7.

Comment

- Quiescence is not identifiable from endpoint counts: the posterior stays near the prior whatever the truth.
- Quiescence, lifetime and self-renewal lie on one ridge (the growth rate).
- Fixing it to 0 gives a constant small shift (longer lifetimes absorb it), independent of the true value, undetected by the PPC.
- Consequence: fix A persist from independent evidence (Ki67) or report lifetimes and self-renewal conditional on it; measuring A's division time also breaks the confound. Q4 (meaning of "inattivazione") open.

Leads to. Part III closes the baseline analysis: GW is the target, BH needs lifetime information. Part IV introduces the errors of the real data one at a time, each compared with these baselines.

Part IV

Misspecified biology and assay

Motivation. The real CRC0322 data will differ from the model in ways Parts I–III did not consider. This part introduces them one at a time, in the order of the experiment: what is seeded (founder type and number, E1a), what dies (E1b, E1c), how cells are labelled (mis-typing, E2a/b) and whether a third population exists (E2c). Each error is measured as a shift from the well-specified control of the same sweep (the BH control is biased, §9) and checked for detectability. Most errors are run in both simulators because the first BH result (§16) suggested that the free time scale hides errors; GW then shows the same error without that freedom. The hidden-type experiment ends by reusing the tail-truth lesson of X5. Section 26 collects detectability across the part.

15 E1a — Wrong founder type (organoid death as initial conditions; BH and GW)

Why this next. All runs so far assume each organoid starts from one A cell. Organoid loss at seeding is assumed (from the biology) to change only which clones start — their founder type and number — not the dynamics. The first test: B-founded organoids.

Question

If some organoids are founded by B cells (fraction $1 - f$) but the model assumes A founders, are the dynamics (A, B rows) biased? Does inferring the founder mix help?

Setup

Simulator	Bellman–Harris (shape 8) and Galton–Watson
Law (free)	`reversible_ab`: BH 9, GW 7; inferred arm adds `founder.A` (Dirichlet(1,1))
Data model	founder A with prob. $f \in \{1, 0.9, 0.7, 0.5\}$, else B; dynamics $\theta^\star$
Inference model	fixed arm: founder = A (wrong for $f < 1$); inferred arm: founder share free
Design	real-like: BH 144/240 h, GW generations 6/10; 165/101 organoids
Inference / budget	NPE amortised, 5k
Seeds / runs	5 / 80 (4 f × 2 arms × 2 simulators)
Control	$f = 1$ of the same arm (same trained network per seed: paired shifts)
Folder	`results/e1a/founder_type/{fixed,inferred}_{bh,gw}`

Result

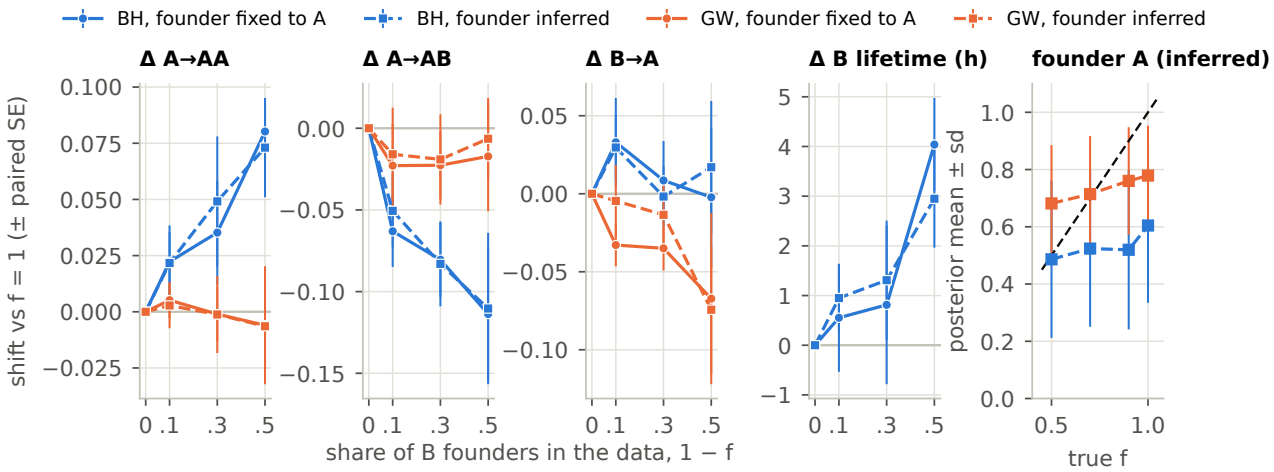


Figure 19: Shift of posterior means vs the $f = 1$ control (± paired SE over 5 seeds); blue BH, orange GW; solid = founder fixed to A, dashed = founder inferred. Right: inferred founder A share vs truth.

	BH, founder fixed				GW, founder fixed			
	control	$f = 0.9$	0.7	0.5	control	0.9	0.7	0.5
A→AA (0.55)	0.32 [0/5]	+0.02	+0.04	+0.08 ± .01	0.52 [3/5]	+0.01	−0.00	−0.01
A→AB (0.05)	0.23 [3/5]	−0.06	−0.08	−0.11± .04	0.12 [3/5]	−0.02	−0.02	−0.02
A→B (0.05)	0.07	+0.01	+0.01	+0.02 ± .006	0.035	+0.00	+0.01	+0.02 ± .006
B death (0.10)	0.17	−0.02	+0.00	−0.02	0.10	+0.01	+0.02	+0.03 ± .014
B→A (0.80)	0.63 [5/5]	+0.03	+0.01	−0.00	0.76 [5/5]	−0.03	−0.04	−0.07± .06
τ_B (30 h)	20.2	+0.6	+0.8	+4.0 ± 0.9			—	
MAE (probabilities)	.109	.090	.095	.089	.033	.033	.036	.047
PPC	.008	.008	0	0	.004	.008	.017	.017
support mass	n/r	n/r	n/r	n/r	0.67	0.69	0.69	0.72

Founder inferred: `founder.A` posterior mean (contraction) for $f = 1$ / 0.9 / 0.7 / 0.5:
BH 0.60 / 0.52 / 0.52 / 0.49 (0.06–0.12) GW 0.78 / 0.76 / 0.71 / 0.68 (0.49–0.62); GW support mass 0.39–0.59

Comment

- **Small shifts, PPC blind** in both simulators. BH: A→AB −0.11, A→AA +0.08 (towards the truth, by accident of the control's bias), τ_B +4 h at $f = 0.5$; all < 1 posterior sd. GW: shifts ≤ 0.07 (B→A, not resolved: ±0.06); A→B +0.02 and B death +0.03 resolved but ≈0.5 sd.
- **GW well-specified baseline is good**: MAE 0.033, contraction 0.89, A→AA 0.52 — the BH bias is specific to BH (ridge).
- **Founder mix**: not identifiable in BH (stays at prior 0.5); partly in GW (monotone but compressed: 0.68 at $f = 0.5$, 0.78 at $f = 1$). Inferring it removes none of the shifts in either simulator and costs τ_B (−3.9 h, BH).
- Plausible reason: B founders switch to A quickly (B→A = 0.8), so B- and A-founded clones look alike by day 6.
- Consequence: keep single-A founders unless measured at seeding; if measured, fix to the measured value. $f < 0.5$ not tested.

Leads to. Founder type is silent and small. The other initial condition is the number of founder cells (§16).

16 E1a — Wrong number of founder cells (BH)

Why this next. Wrong founder type was harmless (§15). Wells may also contain multi-cell seeds, which make clones larger from the start — an effect that could be mistaken for faster growth.

Question

If some organoids start from 2–3 A cells (P(1 cell) = p_1) but the model assumes single-cell seeding, are the dynamics biased? Can the founder size be inferred?

Setup

Simulator	Bellman–Harris, shape 8
Law (free)	`reversible_ab` (9); inferred arm adds `founder_size` (sizes 1–3, Dirichlet(8,1,1))
Data model	founder size $\{1, 2, 3\}$ with $p_1 \in \{1, 0.85, 0.6\}$ ($\{.85, .12, .03\}$, $\{.6, .3, .1\}$), all A
Inference model	sizes [1] (wrong for $p_1 < 1$) or [1,2,3] inferred
Design	real-like 144/240 h × 165/101
Inference / budget	NPE amortised, 5k
Seeds / runs	5 / 30
Folder	`results/e1a/founder_number_bh/` (GW: §17)

Result

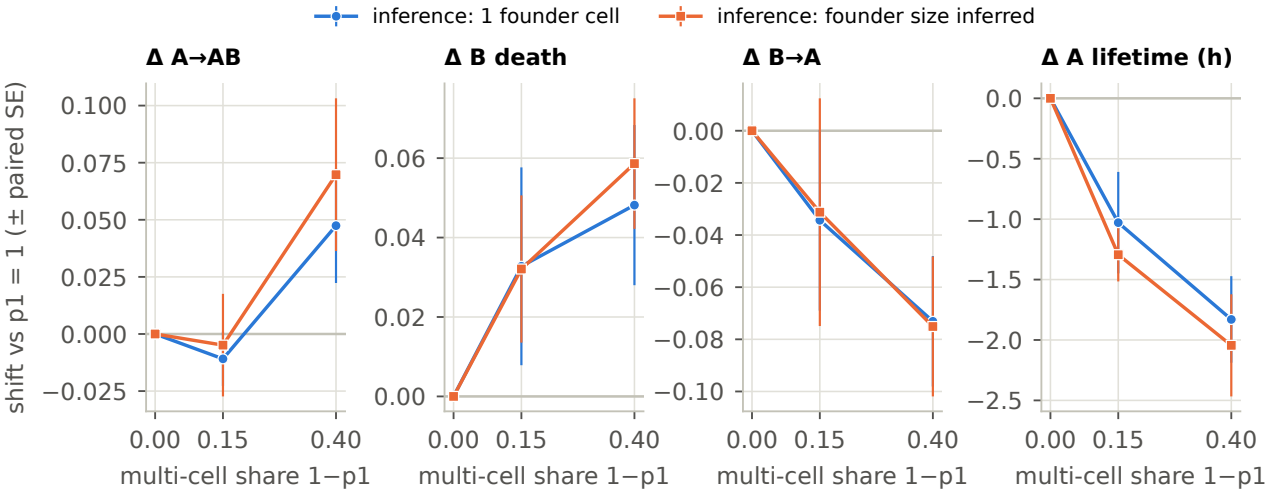


Figure 20: Shift vs the $p_1 = 1$ control of the same arm (± paired SE).

Arm, p_1	A→AB (.05)	B death (.10)	B→A (.80)	τ_A (20 h)	`founder_size.1`
[1], control	0.23	0.17	0.63 [5/5]	15.9	fixed
[1], 0.85 / 0.6	−0.01 / +0.05	+0.03 / +0.05 ± .02	−0.03 / −0.07 ± .03	−1.0 / −1.8 ± 0.4	fixed
[1,2,3], 0.85 / 0.6	−0.00 / +0.07	+0.03 / +0.06	−0.03 / −0.08	−1.3 / −2.0	0.81 / 0.78 (contr. ≤ 0.11)

Comment

- 40% multi-cell clones shift mostly lifetimes and B's row; coverage unchanged; PPC ≤ 0.008 (silent).
- Founder size stays at its prior (0.78–0.83); inferring it does not remove the shifts.
- Observed fraction rises with multi-cell seeding (0.95 → 0.97–0.98): fewer clones go extinct.
- Consequence: tolerable at $p_1 \geq 0.6$; do not infer founder size; measure it at seeding if needed. No support mass recorded (pre-`43e9e44` code).

Leads to. The shifts landed mostly on the lifetimes, and no check fired: does the free time scale absorb the error (as it absorbs self-renewal on the ridge, §9)? The same sweep in GW (§17) answers this.

17 E1a — Wrong number of founder cells (GW)

Why this next. In BH (§16), 40% multi-cell clones moved τ_A by −1.8 h and B→A by only −0.07, silently. GW has no lifetimes to absorb the extra growth.

Question

As §16 for Galton–Watson: are the dynamics biased by multi-cell founders, can the founder size be inferred, and does the answer depend on the time model?

Setup

Simulator	Galton–Watson
Law (free)	`reversible_ab` (7); inferred arm adds `founder_size` (sizes 1–3, Dirichlet(8,1,1))
Data model	founder size $\{1, 2, 3\}$, $p_1 \in \{1, 0.85, 0.6\}$, all A
Design	real-like generations 6/10 × 165/101
Inference / budget	NPE amortised, 5k (one network per arm and seed: paired shifts)
Seeds / runs	5 / 30
Folder	`results/e1a/founder_number_gw/`

Result

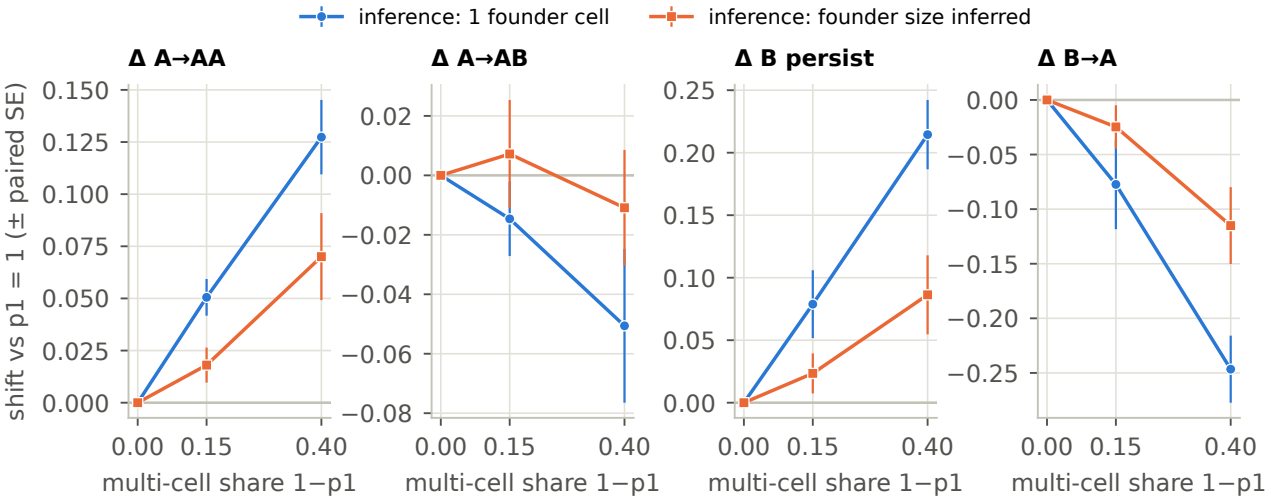


Figure 21: E1a founder number, GW: shift vs the $p_1 = 1$ control of the same arm (± paired SE).

Arm, p_1	A→AA (.55)	B persist (.10)	B→A (.80)	`founder_size.1`	PPC	support (min)
[1], control	0.52 [3/5]	0.13	0.76 [5/5]	fixed	0.004	0.67 (0.44)
[1], 0.85	+0.05 ± .01	+0.08 ± .03	−0.08 [4/5]	fixed	0.042	0.56 (0.21)
[1], 0.6	**+0.13 ± .02** [0/5]	**+0.21 ± .03** [0/5]	**−0.25 ± .03** [0/5]	fixed	**0.117**	**0.16** (0.015)
[1,2,3], control	0.46 [3/5]	0.09	0.83 [5/5]	0.85	0.000	0.55 (0.52)
[1,2,3], 0.85 / 0.6	+0.02 / +0.07 ± .02	+0.02 / +0.09	−0.02 / −0.12 ± .04 [5/5]	0.80 / 0.71	≤ 0.004	0.53 / 0.56

Comment

- **A large, partly detectable error in GW**: 40% multi-cell clones ignored inflate self-renewal (+0.13) and lower B→A (−0.25), all 0/5 covered; the fit is flagged (PPC 0.12, support 0.16). At $p_1 = 0.85$ the shifts are a third of that and the checks borderline.
- **Contrast with BH** (§16): there the same error moved B→A by only −0.07 and was silent — the lifetimes absorb part of the extra growth. A free time scale hides the error.
- **Inferring the founder size helps in GW**: misfit removed (PPC ≤ 0.004), shifts halved and covered, `founder_size.1` tracks p_1 (0.85 / 0.80 / 0.71, contraction ≈ 0.2). Cost on single-cell data: A→AA −0.06 ± 0.01, mean contraction 0.89 → 0.73.
- Consequence: record at seeding whether wells contain multi-cell seeds; if ≳ 15% of organoids start from several cells, fix the measured distribution or infer it. $p_1 < 0.6$ not tested.

Leads to. Founders are one route by which organoid death enters; per-type cell death also changes which organoids are observed at all (E1b, §18).

18 E1b — Death rates and the observed-organoid fraction (BH and GW)

Why this next. E1a (§15–§17) modelled organoid loss through initial conditions. The real data report only organoids that are still detected, and cell death controls how many go extinct: if the number seeded is known, the observed fraction is an extra datum.

Question

How strongly do A death and B death change the fraction of seeded organoids still detected at the endpoints? Is it an informative extra datum?

Setup

Simulator	BH (simulated, rejection loop) and GW (exact $1 - P$(no captured cell))
Law	`reversible_ab`; A death ∈ {0, .03, .1, .2, .3} × B death ∈ {0, .1, .3, .5}, other entries of each row rescaled in proportion
Data vs inference	no inference used (token budget: 256 simulations, 5 epochs; posteriors meaningless)
Design	real-like: BH 144/240 h, GW generations 6/10; 165/101 recovered organoids
Seeds / runs	1 dataset per cell / 40 (20 BH + 20 GW)
Folder	`results/e1b/bh`, `results/e1b/gw`

Result

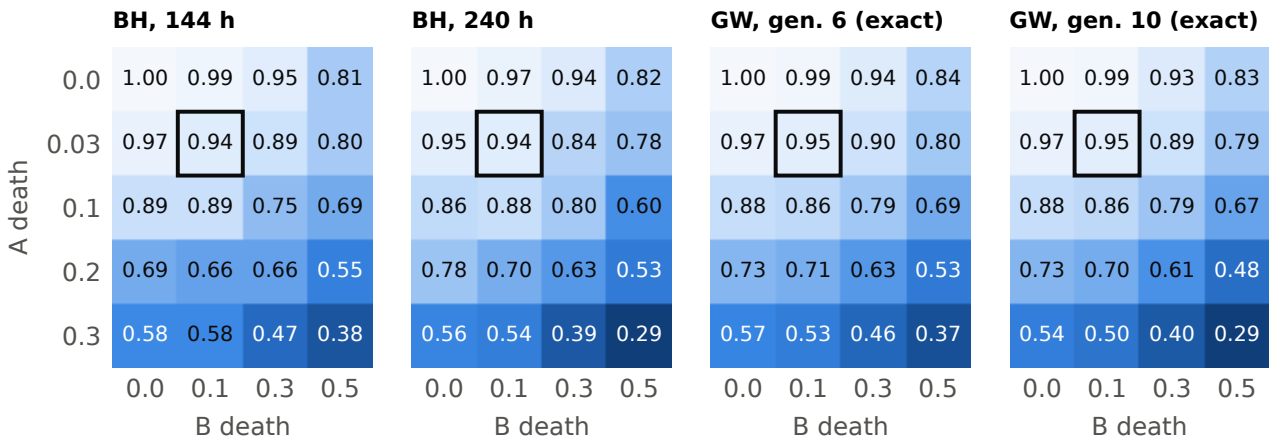


Figure 22: Observed-organoid fraction (detected/seeded) per (A death, B death); box = reference truth. BH: binomial SE ≈ 0.02–0.03 (100–530 seeded clones); GW exact.

(A death, B death)	BH 144 h	BH 240 h	GW gen. 6	GW gen. 10
(0, 0)	1.00	1.00	1.000	1.000
(0.03, 0.1) reference	**0.94**	**0.94**	**0.953**	**0.952**
(0.3, 0)	0.58	0.56	0.565	0.543
(0, 0.5)	0.81	0.82	0.840	0.832
(0.1, 0.3)	0.75	0.80	0.793	0.785
(0.3, 0.5)	0.38	0.29	0.366	0.290

Comment

- A death dominates (0 → 0.3: 1.00 → 0.54–0.58); B death 0 → 0.5 only to ≈ 0.82.
- Losses happen early (day 6 ≈ day 10), except high B death with A death ≥ 0.1.
- Different (A, B) death pairs give the same fraction — confirmed by the exact GW values ((0, 0.5) and (0.1, 0.3): 0.84 vs 0.79; (0.03, 0.5): 0.80): one constraint on a ridge, not two measurements.
- GW exact and BH simulated agree on average to 0.02 (max 0.07, both signs; within ≈2 binomial SE of BH) at the same days/generations.
- Consequence: record the number of seeded organoids per well (Q5). A measured fraction well below 0.9 would indicate A death ≳ 0.1. Not yet tested: using it as a summary (`observation=all_clones`).
- Caveat: death also lowers the growth rate here (rows rescaled).

Leads to. A and B death trade off along one constraint. If death rates must instead be fixed from a guess or an apoptosis marker, do the other parameters compensate (E1c, §19)?

19 E1c — Death rates fixed to wrong values (BH)

Why this next. E1b (§18) showed that different (A, B) death pairs produce the same observed fraction. Fixing deaths is tempting (fewer free parameters), but a wrong value could be compensated elsewhere.

Question

If A and B death are fixed (possibly wrongly) and the rest inferred, do other parameters compensate (flat PPC valley = death not identifiable once fixed)?

Setup

Simulator	Bellman–Harris, shape 8
Law (free)	`reversible_ab` with A death, B death fixed (7 free)
Data vs inference	data always $\theta^\star$ (A death 0.03, B death 0.1); inference fixes A death $\in \{0, 0.03, 0.2\}$, B death $\in \{0, 0.1, 0.4\}$
Design	real-like 144/240 h × 165/101
Inference / budget	NPE amortised, 5k (retrained per setting)
Seeds / runs	5 / 45 (3 × 3 grid; the 4 × 4 grid is the GW sweep, §20)
Control	deaths fixed at truth (0.03 / 0.1)
Folder	`results/e1c/bh/`

Result

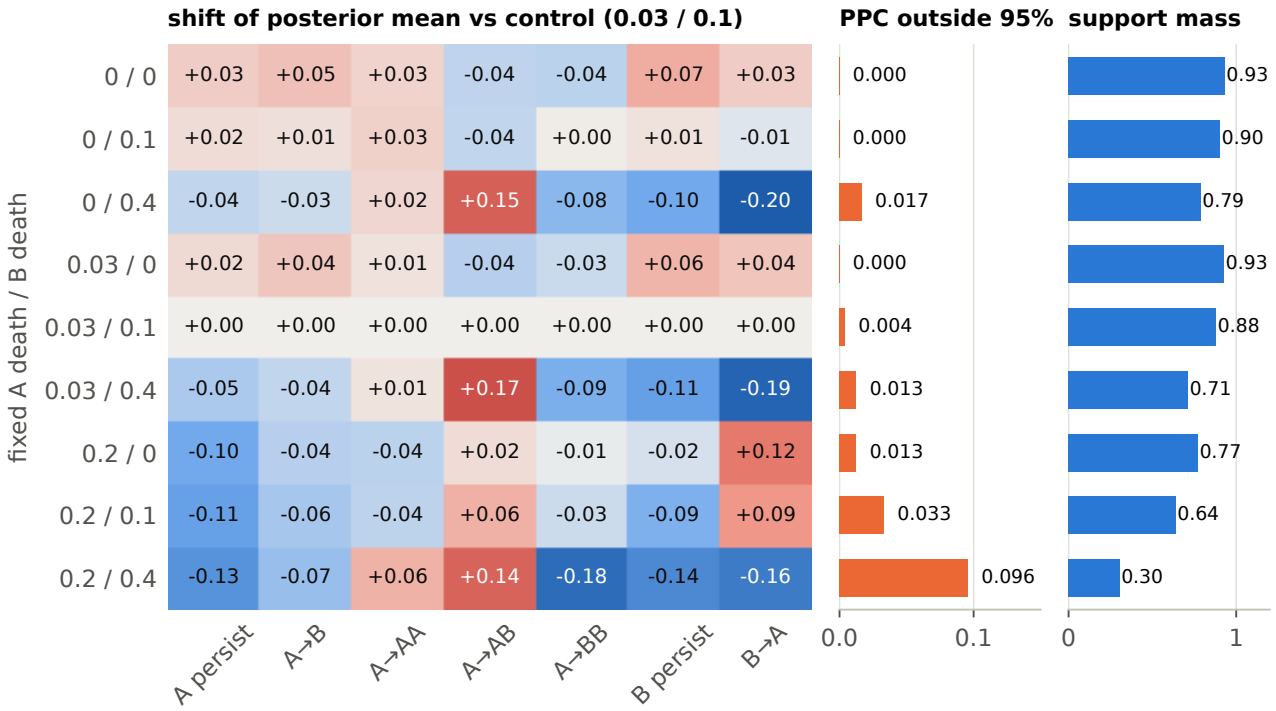


Figure 23: Left: shift of posterior means vs the control (deaths fixed at truth). Right: PPC and posterior support mass (well-specified free-death baseline: support 0.78).

Fixed A / B death	A persist	A→AB	B→A	τ_A	τ_B	PPC	support
0.03 / 0.1 (control)	0.15	0.18	0.67 [5/5]	16.1	18.1	0.004	0.88
0 / 0.4	−0.04	+0.15	−0.20 [0/5]	+0.3	+3.2	0.017	0.79
0.03 / 0.4	−0.05	+0.17	−0.19 [0/5]	−0.3	+2.9	0.013	0.71
0.2 / 0	−0.10	+0.02	+0.12	−2.1	−3.4	0.013	0.77
0.2 / 0.1	−0.11	+0.06	+0.09	−2.7	−2.1	0.033	0.64
0.2 / 0.4	−0.13	+0.14	−0.16 [0/5]	−2.4	+7.9	**0.096**	**0.30**
0 / 0, 0 / 0.1, 0.03 / 0	≤ 0.03	−0.04	≤ 0.04	≤ 0.8	≤ 2.4	0.000	0.90–0.93

Paired SE ≤ 0.03 (probabilities), ≤ 0.35 h (τ_A), ≤ 1.6 h (τ_B).

Comment

- Wrong deaths are absorbed by biologically sensible compensations: B death too high ⇒ less B→A, more A→AB; A death too high ⇒ less quiescence/switching, faster A clock.
- Bias, not information loss: mean contraction *rises* (0.72 → 0.84) while coverage falls (0.76 → 0.38): wrong models are more confident.
- Only the double error is flagged (PPC 0.096, support 0.30; seed 2 support 0.014 ⇒ not an estimate).
- Fixing deaths at the truth gains little vs free deaths (free baseline already constrains them: contraction 0.89 / 0.65).
- Consequence: do not fix death to a guess; keep free or measure (apoptosis marker). C2ST not run; intermediate values (0.1, 0.2) only in the GW sweep (§20).

Leads to. A death too high was absorbed by a faster A clock, again the free time scale (as in §16). The finer GW grid (§20) tests the same errors without it.

20 E1c — Death rates fixed to wrong values (GW, 4 × 4 grid)

Why this next. In BH (§19) only the double error was flagged; A death 0.2 was absorbed silently by a faster clock. Without lifetimes, is a wrong fixed death visible?

Question

As §19, on a finer grid and without the BH ridge: is there a flat PPC valley across fixed death values?

Setup

Simulator	Galton–Watson
Law (free)	`reversible_ab` with A, B death fixed (5 free)
Data vs inference	data always $\theta^\star$; inference fixes A death $\in \{0, .03, .1, .2\}$ × B death $\in \{0, .1, .2, .4\}$
Design	real-like generations 6/10 × 165/101
Inference / budget	NPE amortised, 5k (retrained per setting); bounded SBC (relaunched after the first launch hung; 1 SBC case excluded in 1 run)
Seeds / runs	5 / 80
Control	deaths fixed at truth (0.03 / 0.1)
Folder	`results/e1c/gw/`

Result

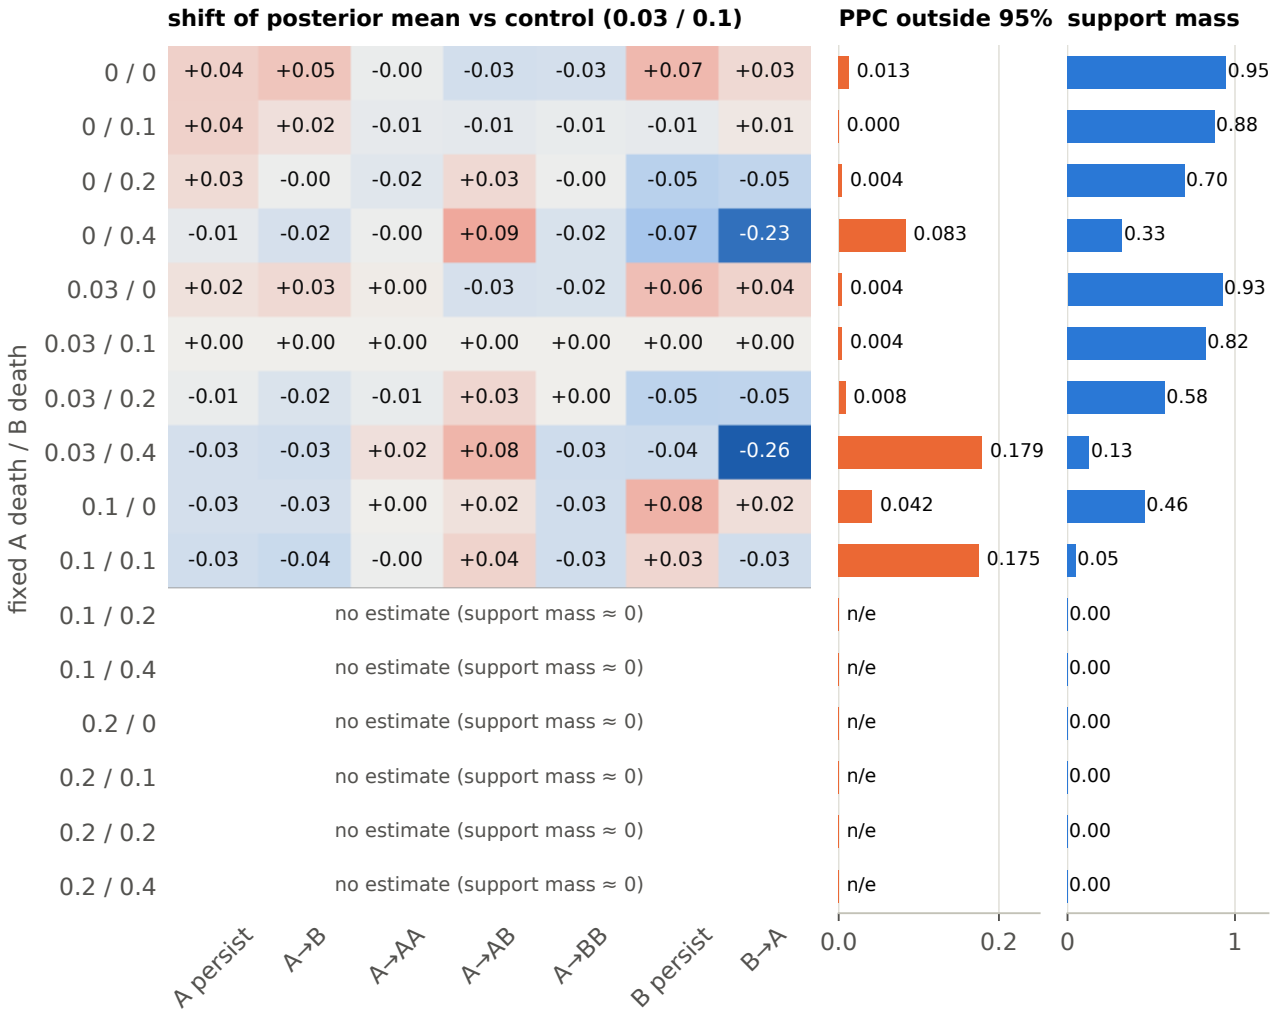


Figure 24: E1c GW. Left: shift of posterior means vs the control (deaths fixed at truth); rows with support mass < 0.02 are not estimates. Right: PPC and posterior support mass.

Fixed A / B death	A persist	A→AB	B persist	B→A	PPC	support
0.03 / 0.1 (control)	0.04	0.10	0.14	0.76 [5/5]	0.004	0.82
0 / 0, 0 / 0.1, 0.03 / 0	$\leq +0.04$	≥ -0.03	$\leq +0.07$	$\leq +0.04$	≤ 0.013	0.88–0.95
0 / 0.2, 0.03 / 0.2	$+0.03$ / -0.01	$+0.03$	-0.05	$-0.05 \pm .02$ [0/5]	≤ 0.008	0.58–0.70
0 / 0.4	-0.01	$+0.09$	-0.07	-0.23 [0/5]	**0.083**	**0.33**
0.03 / 0.4	-0.03 [1/5]	$+0.08$	-0.04	-0.26 [0/5]	**0.179**	**0.13**
0.1 / 0	-0.03 [0/5]	$+0.02$	$+0.08$	$+0.02$	0.042	0.46
0.1 / 0.1	-0.03 [0/5]	$+0.04$	$+0.03$	-0.03	**0.175**	**0.05**
0.1 / 0.2, 0.1 / 0.4	no estimate (support mass ≤ 0.004)				—	**0.00**
0.2 / any	no estimate (17/20 runs without any posterior sample)				—	**0.00**

A→AA never moves ($|\Delta| \leq 0.02$) in an estimable cell. Paired SE ≤ 0.02 for the shifts shown.

Comment

- **The valley is flat downwards, walled upwards for A**: death fixed at 0 fits as well as the truth (shifts ≤ 0.07, PPC ≤ 0.013; small, though up to ≈ 2 posterior sd for GW's narrow posteriors). A death fixed above the truth breaks the fit from 0.1 on (support ≈ 0; at 0.2 no posterior mass in the prior support): GW has one event per generation, so a higher death must be absorbed within A's simplex (A persist and A→B hit 0 first).
- **Contrast with BH** (§19): there A death 0.2 was absorbed silently by a faster A clock (PPC ≤ 0.033, support 0.64–0.77, A persist -0.10, τ_A -2 to -3 h). Whether a wrong fixed death is *detectable* depends on whether the model has a free time scale.
- **B death too high biases B→A in both**: 0.4 $\Rightarrow$ $-0.23/-0.26$ (GW, flagged: support 0.13–0.33) vs $-0.19/-0.20$ (BH, silent). B death 0.2 is silent in GW but already 0/5 covered (B→A -0.05; with B death fixed, B's row has one free coordinate and a narrow posterior).
- Fixing deaths at truth gains little over free deaths (MAE 0.028 vs 0.033). Consequence unchanged: do not fix death rates to guesses. C2ST not run.

Leads to. After seeding and death, the next source of error is the assay: cell types are assigned by clustering, which mislabels cells (E2a/b, §21).

21 E2a / E2b — A↔B mis-typing, known vs ignored (BH)

Why this next. Founders and death (§15–§20) concern the biology. The counts themselves come from marker clustering, with an estimated $\approx 5\%$ double-positive and $\approx 15\%$ double-negative cells assigned to A or B.

Question

Clustering mislabels cells at rate e. If the error is modelled (E2a), do we only lose information? If ignored (E2b), how large is the bias, and is it visible?

Setup

Simulator	Bellman–Harris, shape 8
Law (free)	`reversible_ab` (9)
Data model	symmetric confusion $e \in \{0, .02, .05, .1, .15, .2\}$, capture 1
Inference model	E2a: `observation=noisy` with the same e (network retrained per e); E2b: perfect typing (one network per seed)
Design	real-like 144/240 h × 165/101
Inference / budget	NPE amortised, 5k
Seeds / runs	5 / 60
Control	$e = 0$ (= BH baseline)
Folder	`results/e2/known_bh`, `results/e2/ignored_bh`

Result

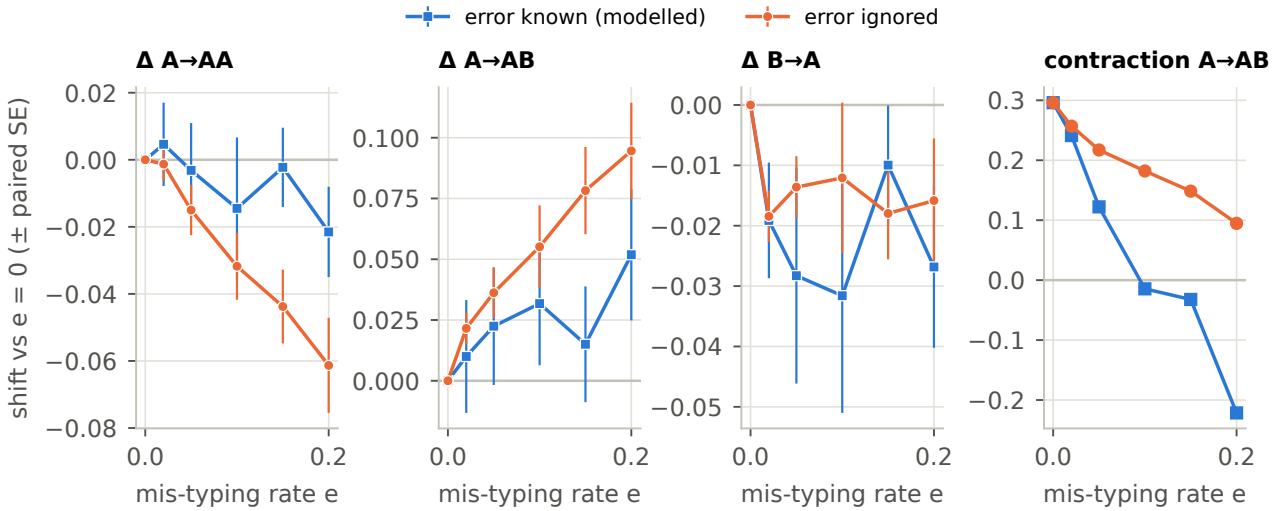


Figure 25: Shift vs $e = 0$ (± paired SE) and contraction of A→AB. Blue: error modelled; orange: error ignored.

e	0.02	0.05	0.1	0.15	0.2	PPC / support
Ignored: A→AB shift	+0.02	+0.04	+0.06	+0.08	+0.09 ± .02	$\leq$ 0.008 / 0.78
Ignored: A→AA shift	−0.00	−0.02	−0.03	−0.04	−0.06± .01	
Ignored: shift / posterior sd (A→AB)	0.2	0.3	0.4	0.6	0.7	
Known: contraction A→AB (0.30 at e=0)	0.24	0.12	−0.01	−0.03	−0.22	$\leq$ 0.008 / n/r
Known: contraction A→AA (0.52)	0.53	0.48	0.43	0.41	0.38	

Comment

- Known error = information loss in the division split; means within 0.05 of control, no trend.
- Ignored error = monotone bias: mislabelling makes clones look mixed, explained by more asymmetric divisions. Coverage of A→AB 3/5 → 0/5 at $e = 0.2$. Invisible to PPC and support mass.
- At the estimated rates (5% ++, 15% −−; Q2) the bias is < 1 sd at this design — tolerable only because posteriors are wide; larger designs (E4) narrow them.
- Consequence: estimate e (2D vs 3D, Q3) and pass it to the model. Asymmetric errors not tested.

Leads to. The bias is tolerable only because BH posteriors are wide; once lifetimes are known (Part III) they narrow to GW-like widths. E2 in GW (§22) shows the tolerance at that precision.

22 E2a / E2b — A↔B mis-typing, known vs ignored (GW)

Why this next. In BH (§21) ignored mis-typing stayed below one posterior sd up to $e = 0.2$, because the posteriors are wide. GW posteriors are $\approx 2.5\times$ narrower: the precision BH would reach with measured division times.

Question

As §21 for Galton–Watson: tolerance to known vs ignored mis-typing when the baseline is unbiased and the posteriors are narrower.

Setup

Simulator	Galton–Watson
Law (free)	`reversible_ab` (7)
Data / inference	symmetric $e \in \{0, .02, .05, .1, .15, .2\}$; modelled (E2a, network retrained per e) or ignored (E2b, one network per seed)
Design	real-like generations 6/10 × 165/101
Inference / budget	NPE amortised, 5k
Seeds / runs	5 / 60
Control	$e = 0$ (= GW baseline)
Folder	`results/e2/known_gw`, `results/e2/ignored_gw`

Result

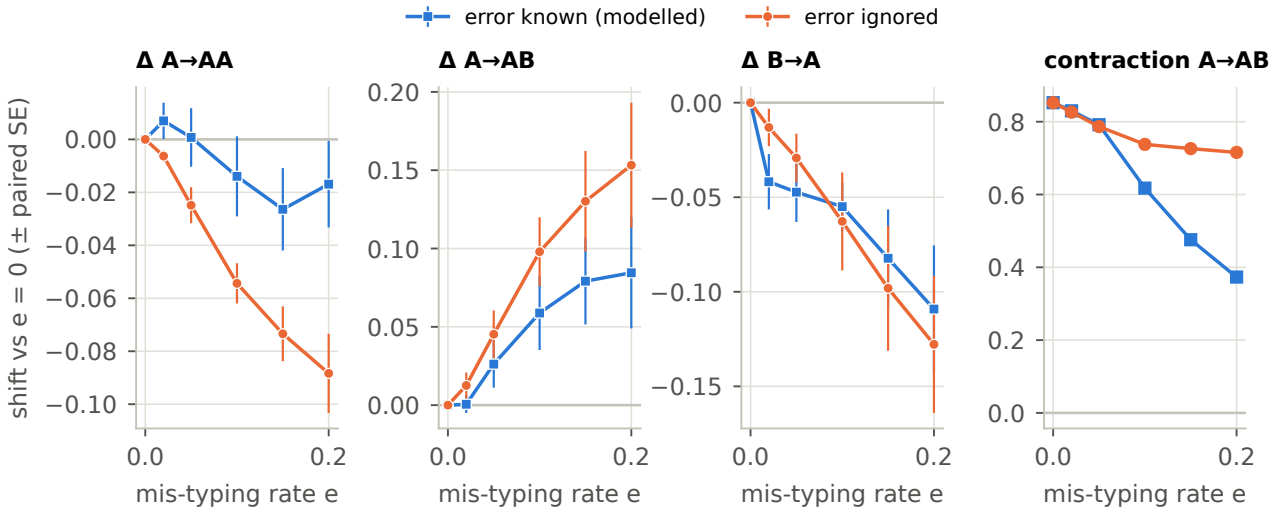


Figure 26: E2 GW: shift vs $e = 0$ (± paired SE) and contraction of A→AB. Blue: error modelled; orange: error ignored.

e	0.02	0.05	0.1	0.15	0.2	BH at 0.2
Ignored: A→AB shift	+0.01	+0.05	+0.10 [0/5]	+0.13	+0.15 ± .04 [0/5]	+0.09
Ignored: A→AA shift	−0.01	−0.02	−0.05 [0/5]	−0.07	−0.09± .01 [0/5]	−0.06
Ignored: B→A shift	−0.01	−0.03	−0.06	−0.10	−0.13± .04 [3/5]	−0.02
Ignored: shift / posterior sd (A→AB)	0.2	0.7	**1.4**	1.8	2.1	0.7
Ignored: PPC / support	.004 / .66	.004 / .64	0 / .59	.012 / .55	.033 / .48	.000 / .78
Known: contraction A→AB (0.85 at e=0)	0.83	0.79	0.62	0.48	0.37	−0.22
Known: B→A shift	−0.04	−0.05	−0.05	−0.08	−0.11± .03 [4/5]	−0.03

Comment

- **Same mechanism as BH**, about 1.5× larger and extending to B's row: mislabelled clones look mixed ⇒ more asymmetric division, less self-renewal, fewer B→A, more B persist (+0.18 at $e = 0.2$).
- **Tolerance is set by the posterior width**: GW posteriors are ≈2.5× narrower (sd of A→AB 0.05 vs 0.12), so the bias exceeds one posterior sd at $e \approx 0.1$ (0/5 covered). Tolerable $e \approx 0.05$ in GW vs ≈ 0.2 in BH.
- **Modelled error = information loss**: A→AB contraction 0.85 → 0.37; means drift towards the prior means (B→A −0.11, A→AB +0.08 at $e = 0.2$) with coverage kept — expected for a wider posterior under a uniform Dirichlet, not a bias.
- **Not detected**: PPC ≤ 0.033; support mass falls 0.67 → 0.48 (min 0.10) — a trend across the sweep, not a single-run warning.
- Consequence: estimate e (Q3) and model it; the BH tolerance will shrink to the GW one once division times are measured. Asymmetric errors not tested.

Leads to. E2 treats the double-negative/positive cells as random mislabelling. They may instead be a real third type C with its own dynamics (E2c, §23).

23 E2c — Hidden third type ignored (GW)

Why this next. E2 (§21, §22) treated mislabelled cells as noise. If they are a real population C that divides and differentiates, ignoring it is a model error, not a labelling error.

Question

Data contain a third type C (double negative/positive) read as A with prob. q and as B with $1 - q$. If the model ignores C, are the A and B rows still recovered, and does the PPC flag the two-type model?

Setup

Simulator	Galton–Watson
Data model	`three_type`: A = {death .03, persist .02, A→C, A→AA, A→AC, A→CC}, C = {death .05, persist .10, C→B .40, C→BC .20, C→CC .25}, B = {death .10, persist .90} (B never returns to A). C fraction ≈ 5% (A→AA 0.905) or ≈ 15% (A→AA 0.79); `three_split` with $q \in \{0, 0.5, 1\}$
Inference model	two-type `reversible_ab`, perfect typing (7 free)
Truth for scored names	A death .03, A persist .02, A→AA .905 / .79, A→B = A→AB = A→BB = 0, B death .10, B persist .90, B→A 0
Design	real-like generations 6/10 × 165/101
Inference / budget	NPE amortised, 5k
Seeds / runs	5 / 30
Control	*none* in this sweep (no two-type data with the same A row); arms with C modelled: §24
Folder	`results/e2c/ignore_c/`

Result

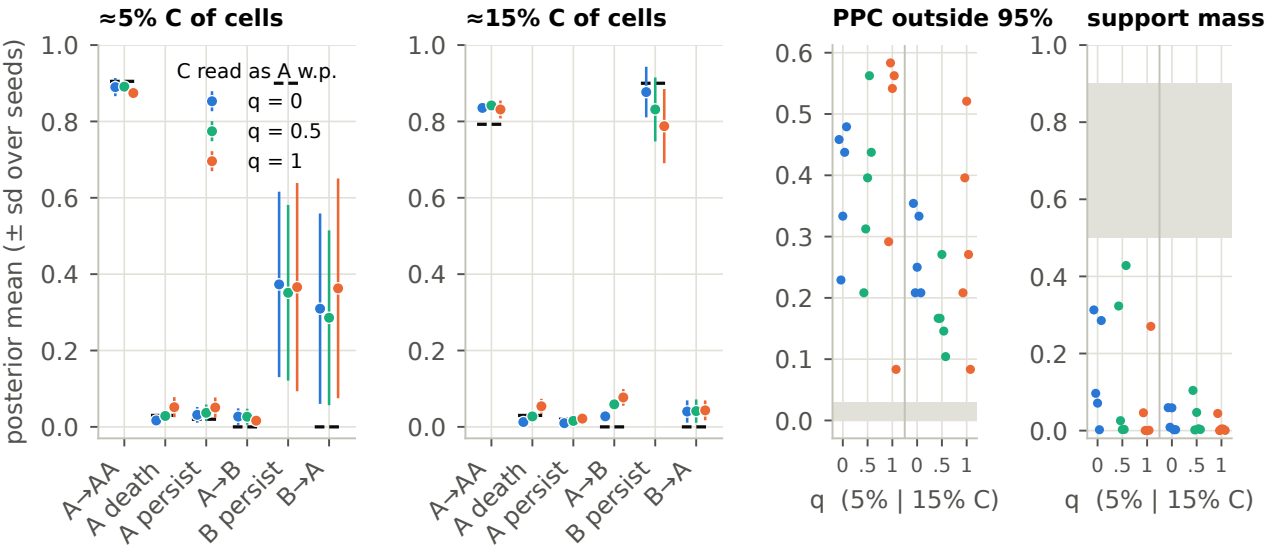


Figure 27: Left: posterior means (± sd over seeds) per q; dashes = truth of the data model. Right: PPC and support mass per run; grey = well-specified ranges.

C share	q	A→AA	A→BB (0)	A→B (0)	B persist (.90)	B→A (0)	PPC	support
≈15%	0	0.84 (.79)	0.10	0.03	0.88	0.04	0.27	0.03
	0.5	0.84	0.05	0.06	0.83	0.04	0.17	0.03
	1	0.83	0.02	0.08	0.79	0.04	0.30	0.01
≈5%	0	0.89 (.905)	0.03	0.03	0.37	0.31	0.39	0.15
	0.5	0.89	0.01	0.03	0.35	0.29	0.38	0.16
	1	0.87	0.01	0.02	0.37	0.36	0.41	0.06

Comment

- **The two-type fit is flagged**: PPC 0.17–0.41 (well specified ≤ 0.03), support mass 0.01–0.16 (well specified 0.5–0.9). Per the reading rules these posteriors are not reliable estimates.
- How C is absorbed: read as B ($q = 0$), A→CC appears as A→BB (0.10 at 15%); read as A ($q = 1$), C→B appears as A→B switching (0.08) and A death rises (0.05). A→AA is overestimated at 15% (+0.05) and slightly under at 5% (−0.02).
- At 5% C, B cells are rare (B only arises via C): B's row stays prior-like (B persist 0.37, contraction 0.25–0.45 for B death).

- **Attribution**: the data truth is far in the prior tail of the inference model (A→AA 0.79–0.905: prior mass above $<0.05\%$; B persist 0.9: 1%), and PPC is *higher* with less C (5% vs 15%). The arms that model C (§24) fail even more although one of them is the data model: the misfit is driven by the prior tail and cannot be attributed to C.
- One run (5%, $q = 1$, seed 2) drew only 74 posterior samples.

Leads to. Ignoring C is flagged, but the misfit cannot yet be attributed to C. Modelling C (§24) should remove it if C is the cause.

24 E2c — Third type modelled: three-arm comparison (GW)

Why this next. The two-type fit to three-type data was flagged (§23), but the PPC rose with *less* C. If the misfit is due to C, modelling it should remove it.

Question

When C is modelled (freely, or with the trajectory constraint C→B only), are A and B recovered, is C's row identifiable without its own marker, and is the misfit of §23 removed?

Setup

Simulator	Galton–Watson
Data model	as §23 (same 30 datasets: $q \in \{0, .5, 1\}$, ≈5% / 15% C)
Inference arms	ignore C: `reversible_ab` (7 free, §23); C free: `three_type` + `three_split`, q known (27 free) = *the data model*; C trajectory: `three_type_traj` (C never produces A; pending Q1) + `three_split` (23 free)
Design	real-like generations 6/10 × 165/101
Inference / budget	NPE amortised, 5k
Seeds / runs	5 / 30 per arm (90)
Folder	`results/e2c/{ignore_c,model_c_free,model_c_traj}`

Result

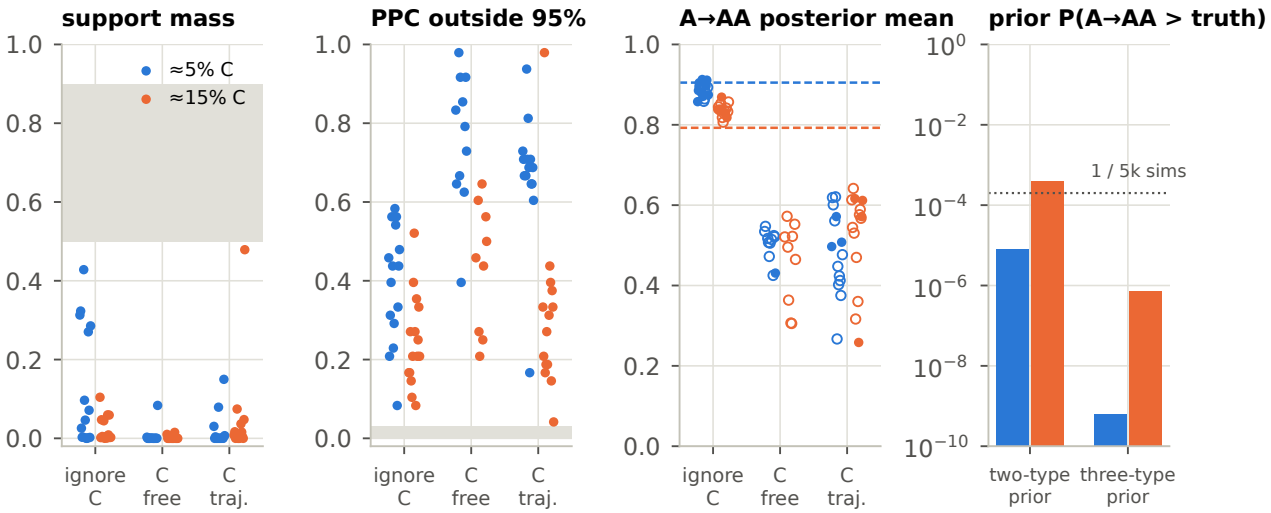


Figure 28: E2c, three inference arms on the same data (one point per run). Grey bands: well-specified ranges. A→AA: open circles = support mass < 0.02 (not estimates); dashed = truth (0.905 at 5% C, 0.79 at 15%). Right: prior probability of A→AA above the truth under each inference model (uniform Dirichlet; dotted = one in 5 000 training simulations).

Arm	C	runs with samples	support: median (max)	PPC (group means)	A→AA (truth)
ignore C	5%	15/15	0.05 (0.43)	0.38–0.41	0.87–0.89 (.905)
	15%	15/15	0.004 (0.10)	0.17–0.30	0.83–0.84 (.79)
C free (well specified)	5%	11/15	0.000 (0.08)	0.63–0.84	*0.48–0.52*
	15%	9/15	0.000 (0.02)	0.37–0.50	*0.33–0.53*
C trajectory	5%	14/15	0.000 (0.15)	0.60–0.75	*0.39–0.60*
	15%	14/15	0.008 (0.48)	0.26–0.38	*0.35–0.59*

Italics: not estimates (support ≈ 0). SBC coverage error 0.024–0.029 in every arm. Prior P(A→AA > truth) at 5% / 15% C: two-type (6 outcomes) $8 \cdot 10^{-6}$ / $4 \cdot 10^{-4}$; three-type (10 outcomes) $6 \cdot 10^{-10}$ / $8 \cdot 10^{-7}$. P(B persist > 0.9): 0.01 (two-type) vs 10^{-9} (three-type).

Comment

- **Modelling C does not remove the misfit — it makes it worse**: support ≈ 0 in almost every run (10/30 free and 2/30 trajectory runs drew no posterior sample at all), PPC 0.26–0.84. No A/B bias and nothing about C (e.g. C→B 0.01–0.12 vs 0.40) can be read from these arms.
- **The failure is a prior-tail failure, not misspecification**: the free arm *is* the data model. Under a 10-outcome uniform Dirichlet the true A→AA (0.79–0.905) and B persist (0.9) have prior probability 10^{-6}–10^{-10}: no training simulation resembles the data, and amortised NPE extrapolates outside the prior support. SBC is fine because it averages over the prior. The two-type model fits *better* than the true model because its prior has more mass near the truth.
- **Hence the ignore-C flag (PPC 0.17–0.41) cannot be attributed to C**; also within ignore-C the PPC is higher with *less* C. Whether the PPC detects a hidden type is still open. The trajectory constraint helps slightly (fewer empty posteriors; one run with support 0.48).
- Follow-up done: the rerun with truths drawn from the three-type prior (§25) confirms that this failure was the tail truth, not the method.

Leads to. The true model fails at a truth with prior probability 10^{-6}–10^{-10}: the same pattern as E0 (§3) and the BH ridge, which X5 (§4, §10) traced to truth position. The rerun with truths drawn from the prior (§25) applies the same remedy.

25 E2c rerun — hidden third type, truths from the prior (GW)

Why this next. The original E2c (§23, §24) could not separate a C effect from a tail-truth failure. X5d (§10) showed that inference is calibrated for truths in the prior bulk; drawing the three-type truths from the prior removes the tail problem and makes both E2c questions answerable.

Question

The original E2c (§23, §24) used a truth deep in the three-type Dirichlet tail, so even the true model failed. With truths in the prior bulk: (1) is ignoring C detected, and (2) does modelling C (free, or trajectory-constrained) recover the A/B rows?

Setup

Simulator	Galton–Watson
Data model	`three_type`, 10 truths *drawn from the three-type prior* (`truth.source=prior`, `data_seed` = seed 0–9); `three_split` with $q \in \{0, 0.5, 1\}$ (C read as A with prob. q). Same 30 datasets in every arm
Prior-drawn truths	generic biology: C is 21–41% of cells at generation 10; B divides (B-division mass 0.37–0.84); C makes A (C→A mass 0.19–0.58)
Inference arms	ignore C: `reversible_ab` (7 free; also forbids B divisions); C free: `three_type` + `three_split`, q known (27 free) = *the data model*; C trajectory: `three_type_traj` (C never makes A; 23 free). Both constrained arms are misspecified for *every* truth, ignore C doubly so
Scored	the 9 A/B parameters present in every arm; "flagged" = support mass < 0.1 or PPC > 0.05; "usable" = at least 100 posterior samples within the 60 s sampling cap; prior-mean baseline = $\lvert\text{truth} - 0.1\rvert$ (10-outcome Dirichlet mean)
Design	real-like generations 6/10 × 165/101
Inference / budget	NPE amortised, 5k
Seeds / runs	10 truths × 3 q × 3 arms = 90
Folder	`results/e2c_prior/{ignore_c,model_c_free,model_c_traj}` (`e2c_prior_runs.csv`, `e2c_prior_groups.csv`)

Result

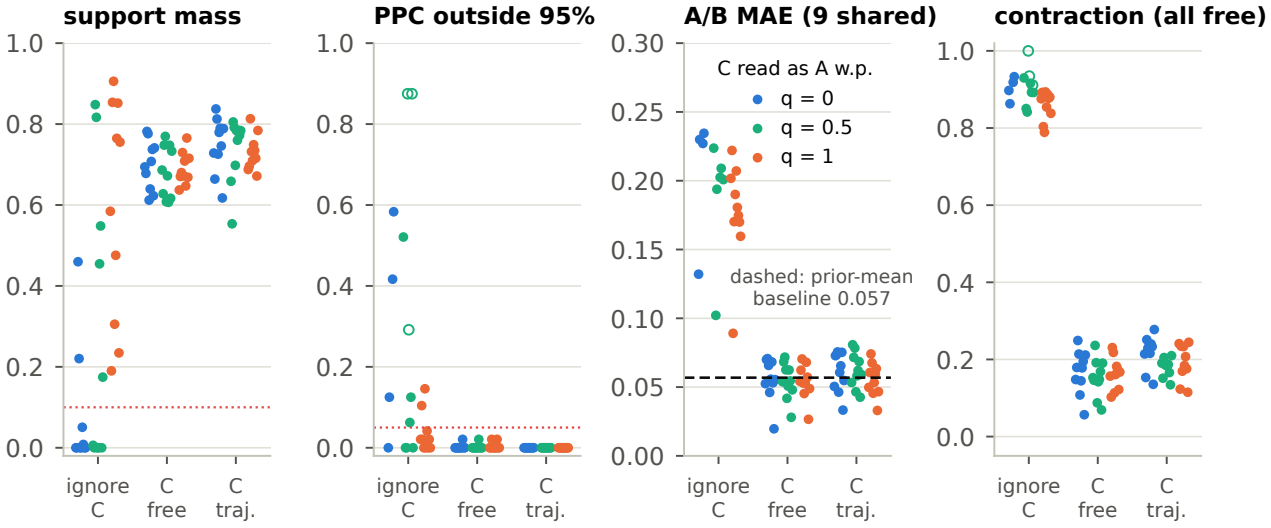


Figure 29: E2c rerun, one point per run (10 prior-drawn truths per q). Dotted red: flag thresholds (support < 0.1, PPC > 0.05). Open circles: fewer than 100 posterior samples (not estimates; A/B MAE not computed). A/B MAE over the 9 shared A/B parameters; dashed = prior-mean baseline. Contraction over all free parameters of each arm.

Arm	q	usable	flagged	support (med.)	PPC	SBC err.	A/B MAE	A/B cov.	contr.
ignore C	0	4/10	**9/10**	0.00	0.28	0.029	0.206	0.11	0.90
	0.5	6/10	**7/10**	0.09	0.31	0.029	0.189	0.20	0.91
	1	10/10	**2/10**	0.67	0.04	0.029	0.177	0.42	0.86
C free (well specified)	0–1	30/30	0	0.68–0.70	0.00	0.026–0.027	0.055	0.93	0.16
C trajectory	0–1	30/30	0	0.72–0.78	0.00	0.024–0.025	0.059	0.92	0.20
Prior-mean baseline (posterior mean replaced by the prior mean 0.1)							0.057		

Model-C rows pool the three q (per q: MAE 0.054–0.056 free, 0.055–0.062 trajectory; contraction 0.15–0.17 and 0.18–0.22). PPC and contraction for ignore C include runs with fewer than 100 samples. Numbers: `tables/e2c_prior.csv`.

Comment

- **Ignoring C is detected when C looks like B, missed when C looks like A.** At $q = 0$ 9/10 runs are flagged (support ≈ 0, PPC up to 0.58; 6/10 runs give no usable posterior). At $q = 1$ only 2/10 are flagged, while the A/B rows are still wrong (MAE 0.18, 90% coverage 0.42) with high contraction (0.86): *confident, silent bias.*
- **Modelling C is calibrated but not informative.** Both model-C arms pass every check (PPC ≈ 0, SBC 0.024–0.027, A/B coverage 0.92–0.93), but contraction is 0.16–0.20 and the A/B error equals the prior-mean baseline (0.055 / 0.059 vs 0.057). With 23–27 free parameters and two endpoints of two labels the data barely move the posterior: correct, but it is the prior speaking.
- **The trajectory constraint is wrong for every truth and nothing detects it** (C→A mass up to 0.58; PPC 0, support 0.72–0.78): the data are too weak to contradict it.
- **Resolves §24**: the "true model fails" result there came from the tail truth, not the method. In the prior bulk the well-specified model is calibrated, but a three-type model is not identifiable from two-label endpoint counts at this design.
- Caveats: 10 truths, NPE 5k, one design. Prior truths have much more C (21–41%) than the biological scenario (5–15%) and B divisions, unlike the original truth; ignore C also forbids B divisions, so its failures are not due to C alone. The A/B "truth" for the two-type arm is the three-type value of the same outcome, which a two-type row cannot match exactly (part of its error is projection).
- Consequence: if a hidden C is suspected, modelling it needs constraints that are actually true (from trajectory data), a C-specific measurement, or much richer designs. Ignoring it is risky mainly when C is read as A (q near 1).

Leads to. This completes the one-at-a-time errors. Section 26 collects which of them the fit itself would reveal; X4 (§28) would combine them.

26 Synthesis — which errors are detectable (BH and GW)

Why this next. Every experiment of this part asked whether the PPC or the support mass flags the error, and a pattern recurred: errors that BH absorbs silently through its free time scale are flagged in GW (§17, §20). This section puts the misspecified settings of both simulators on one scale. It contains no new runs.

Question

Across the misspecification experiments — Part IV and, for BH, the timing errors E3a–E3c of Part III — how large is each shift relative to the control's posterior sd, and would the fit itself have revealed it?

Result

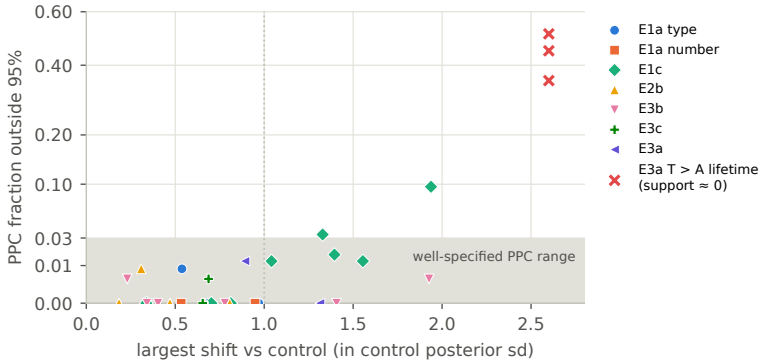


Figure 30: Detectability of BH misspecifications. Each point is one misspecified setting (mean of 5 seeds): largest shift of a posterior mean from its paired well-specified control, in units of the control's posterior sd, against the PPC. Grey band = PPC range of well-specified runs. Crosses: global lifetime $T > \tau_A$ (support mass ≈ 0).

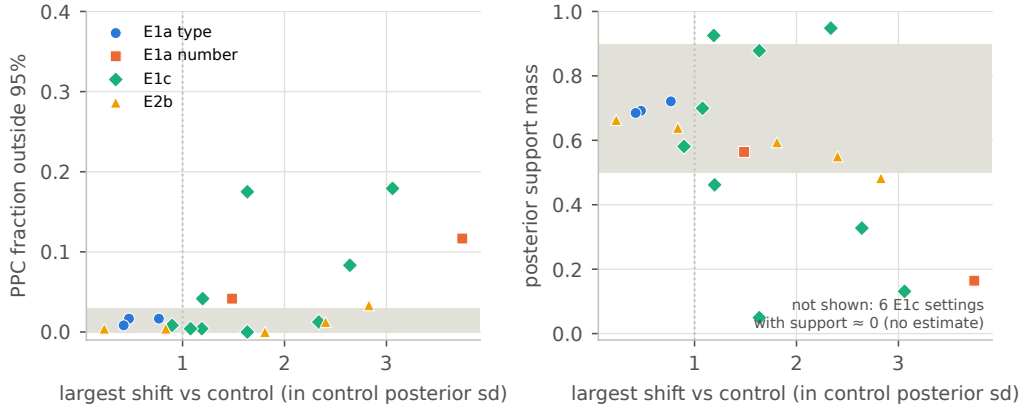


Figure 31: Detectability of GW misspecifications (E1a type and number, E1c, E2b): largest shift from the paired control in control posterior sd, against the PPC (left) and the posterior support mass (right). Grey bands = well-specified ranges. GW flags more errors than BH (no free time scale), but several shifts of 1–3 sd remain inside the well-specified bands.

Comment

- **BH**: the PPC stays ≤ 0.033 for every error except both deaths fixed wrong (0.096, §19) and a global lifetime longer than A's (0.35–0.51, §12).
- **GW**, without lifetimes, flags more: A death fixed too high (support ≈ 0, §20), B death fixed at 4× truth (support 0.13–0.33), 40% multi-cell founders (PPC 0.12, support 0.16, §17).
- **Still silent in GW**: wrong founder type (§15), B death at 2× truth (B→A 0/5 covered), ignored mis-typing up to 2 posterior sd (§22); and a hidden type read as A (§25).
- SBC never detects misspecification. Support mass and PPC vary between well-specified GW runs (support 0.07–0.89, PPC up to 0.10), so single borderline values are not proof.

Leads to. Errors were tested one at a time on 5 seeds; Part V lists the work that would turn them into an error budget for the real-data fit.

Part V
Planned experiments

Motivation. All sweeps of E0–E4, the E2c rerun and X5 are finished (Parts I–IV). The experiments below were never run; no results exist for them. Each is listed with the earlier result that motivates it.

27 X2 — Seed spread: dataset vs training

Why this next. Every comparison in Parts I–IV rests on 5 seeds, each changing both the dataset and the trained network (reading rule 4, §2); single well-specified GW runs already vary widely in support mass and PPC (§7).

Question

How much of the seed-to-seed spread comes from the dataset and how much from network training?

Setup

Configuration	final configuration from E1–E4 (to be chosen)
Arms	training spread: 1 dataset × 5 training seeds; dataset spread: 5 `data_seed` × 1 training seed
Runs	5 per arm (planned, not started)
Folder	`results/x2/training`, `results/x2/data`

Result pending — planned

Will show the posterior mean and sd per parameter in both arms. If the training spread is comparable to the dataset spread, single-run results need ensembles or larger budgets.

28 Other planned work (X1, X3, X4, R, P0, follow-ups)

Why this next. X1: every result uses amortised NPE (SNPE only in E0, §3). X3: the GW/BH contrast recurs throughout (§9, §17, §20, §22), but the two simulators do not share a time unit. X4: Part IV tested errors one at a time. R: the end goal. P0: BH cost was never measured on an idle machine (Appendix A). Follow-ups: the BH ridge at realistic sizes with a well-trained network (§8.1), asymmetric mis-typing (§21), `two_type` at large size (§5), the observed fraction as a summary (§18).

Question

- **X1 methods**: NPE vs NLE vs NRE (and SNPE) at the baseline and the worst realistic error.
- **X3 model**: can a GW fit describe BH data? (needs a GW↔hours mapping; code change).
- **X4 combined errors**: realistic combination (organoid death + mis-typing + C + lifetime CV), real-like design, `budget=full`: the error budget to quote with the real-data fit.
- **R real data**: CRC0322, day 6 (165 organoids) and day 10 (101); choose 2D counts vs 3D estimates (Q3); final model, SBC, fit, PPC.
- **P0 BH timing**: seconds per BH run at `dev`/`smoke` budgets on an idle machine.
- **BH ridge at realistic sizes with a well-trained network (decisive follow-up)**: the E4 budget check (§8.1) moved the ridge at 4096 clones per endpoint with 20k simulations. Repeat with 20k (or more) simulations and a higher epoch cap at the real-like design and at 512 per endpoint (days 4/6/10): if A→AA and τ_A move there too, the BH bias is a training limit at realistic sizes; if not, lifetime information remains necessary.
- **Follow-ups of finished sweeps**: asymmetric mis-typing (E2); E4 GW with `law=two_type` at 256–512 clones; observed fraction as a summary (`observation=all_clones`, E1b).

Setup

X1	`+experiment=compare_methods`; NLE/NRE with 50 SBC cases; 3 seeds
X4	real-like design, 5 seeds, `budget=full`
R	`submodules/GW/data/df_6D_withmapping.csv`, `df_10D_withmapping_new.csv`
Blocking	Q1 (C trajectory), Q2 (C proportions), Q3 (2D vs 3D mislabelling), Q4 (meaning of inactivation), Q5 (seeded organoids known?)

Result pending — planned, not started

X1: error/contraction/SBC per method. X3: bias of a GW fit to BH data. X4: total bias vs posterior sd per parameter. R: posterior on CRC0322 with the error budget of X4.

A Device and batch-size benchmark

Why this next. Engineering, not part of the argument. E0's 50k-simulation and SNPE runs (§3) were the slowest of the project, and wall times recorded in run metrics were measured under heavy contention. Before the large E1–E4 sweeps: does a GPU or a larger batch help?

Question

Do GPU training and/or a larger batch (512 vs 128) make runs faster without degrading the posterior?

Setup

Configuration	GW, `two_type`, generations 4/8/12 × 32, NPE 20k, seed 0, SBC 200, PPC 200, 4 threads
Hardware	20-core CPU; NVIDIA RTX A1000 (8 GB, also driving the display)
Varied	device (CPU/GPU) × batch (128/512); SNPE 4 rounds on CPU and GPU
Protocol	one job at a time, contention checked every 15 s (all 6 jobs clean); single measurements
Folder	`results/gpu_bench/` (branch `gpu-bench`, `811d8ca`)

Result

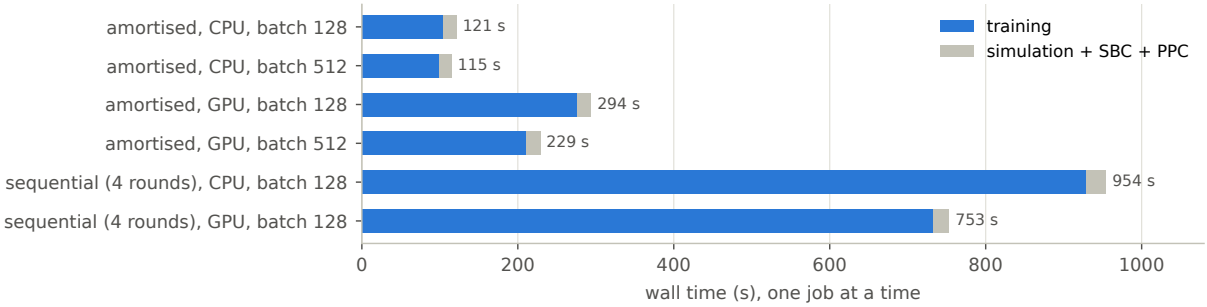


Figure 32: Wall time per job (blue = network training).

Job	wall (s)	s/epoch	epochs	MAE	SBC error
amortised, CPU, 128	121	1.47	71	0.140	0.032
amortised, CPU, 512	115	0.65	151	0.140	0.028
amortised, GPU, 128	294	4.07	68	0.137	0.033
amortised, GPU, 512	229	1.27	166	0.138	0.031
SNPE, CPU, 128	954	4.3	215	0.135	—
SNPE, GPU, 128	753	3.2	226	0.129	—

Comment

- Amortised NPE: GPU 2–2.4× slower (small flows, launch overhead). Batch 512: epochs 2.3× faster but 2× more epochs: −5%.
- SNPE: GPU 21% faster. Quality identical across devices.
- Keep CPU + batch 128; GPU only for SNPE or as extra capacity. These are the only valid timings in this report.

B Reproducibility

Code versions. `manifest.json` in each run records `git_commit` (HEAD at run time; since `6b45d23` also the commit of the loaded code). Nearly all runs have `git_dirty = true`. Several sweeps span several HEAD values because commits were made while they ran; the loaded code is that of the launch commit (verified for GW E1a: identical networks across settings per seed).

Experiment	Launch commit	Command (preset; `hydra.sweep.dir=results/...`)
E0 (§3)	`194e31f`, `b3cbae0`	`+experiment=b_row_identifiability law=two_type|reversible_ab`; `inference.method.rounds=4`; `+experiment=b_row_flat_prior`
X5a (§4)	`d403cac`; bio `7f4a11f`	`+experiment=x5a_truth_position law=two_type|reversible_ab '++law.concentration={},<Jeffreys>,<flat B.to_A>'` biological priors: same with `'++law.concentration=<bio-B>,<bio>'` (folders `results/x5/a_*_bio`) (exact lists in the preset header and `results/x5/a_*/info.md`)
X5c (§5)	`d403cac`	`+experiment=x5c_data_vs_prior`
Design comparison (§6)	`b3cbae0`	`+experiment=design_comparison` (laws two_type, reversible_ab)
E4 GW (§7)	`19a10ea`	`+experiment=e4_design`
	`5cb6018` (1024, 2048)	`+experiment=e4_design design.clones=1024,2048` (folder `results/e4/gw_large`)
	`07fb7fe` (4096)	`+experiment=e4_design design.clones=4096` (folder `results/e4/gw_xxl`)
	`07fb7fe` (20k; docs `503ec72`)	`+experiment=e4_design design.clones=4096 'design.horizons=[4,6,10]' budget.simulations=20000 seed=0,1,2,3,4` (folder `results/e4/gw_xxl_20k`)
E4 BH (§8)	`43e9e44`	`+experiment=e4_design simulator=bellman_harris design.horizons=[240],[144,240],[96,144,240],[144,240,336]`
	`5cb6018` (1024, 2048)	same + `design.clones=1024,2048` (folder `results/e4/bh_large`)
	`07fb7fe` (4096)	same + `design.clones=4096` (folder `results/e4/bh_xxl`)
	`07fb7fe` (20k; docs `503ec72`)	`simulator=bellman_harris design.clones=4096 'design.horizons=[96,144,240]' budget.simulations=20000 seed=0,1,2,3,4` (folder `results/e4/bh_xxl_20k`)
BH ridge (§9)	`19a10ea`	`+experiment=bh_ridge` (+ lifetimes fixed / A concentration)
X5d (§10)	`d403cac`	`+experiment=x5d_prior_average simulator=galton_watson|bellman_harris`
X5b (§11)	`d403cac`	vague: `+experiment=x5b_lifetime_prior` right / wrong: add `'~simulator.model.lifetime_mean_bounds'` and `'+simulator.model.lifetime_mean_bounds={A:[16,25],B:[24,37.5]}'` resp. `{A:[24,37.5],B:[36,56.25]}`
E3a / E3b / E3c (§12–14)	`27b1d61`–`c919397`	`e3a_lifetime_free`, `e3a_lifetime_fixed`, `e3b_desync`, `e3c_persist`
E1a type (§15)	BH `27b1d61`; GW `19a10ea`/`d46b4c2`	`+experiment=e1a_founder_type simulator=...` (fixed; inferred with `'~law.fixed.founder'`)
E1a number BH (§16)	`27b1d61`	`+experiment=e1a_founder_number simulator=bellman_harris`
E1a number GW (§17)	`b0789d3`	`+experiment=e1a_founder_number simulator=galton_watson`
E1b (§18)	BH `43e9e44`; GW `d46b4c2`	`+experiment=e1b_death_detection simulator=...`
E1c BH (§19)	`43e9e44`/`c919397`	`+experiment=e1c_fixed_death simulator=bellman_harris law.fixed.A.death=0,0.03,0.2 law.fixed.B.death=0,0.1,0.4`
E1c GW (§20)	`d46b4c2`	`+experiment=e1c_fixed_death` (4×4 grid; relaunched with bounded SBC)
E2a/E2b BH (§21)	`27b1d61`–`c919397`	`+experiment=e2a_typing_known`, `+experiment=e2b_typing_ignored`
E2a/E2b GW (§22)	`12c0827`	`+experiment=e2a_typing_known|e2b_typing_ignored simulator=galton_watson`
E2c ignore C (§23)	`19a10ea`	`+experiment=e2c_third_type`
E2c C modelled (§24)	free `d46b4c2`; traj. `12c0827`	`+experiment=e2c_third_type law=three_type|three_type_traj observation=three_split`
E2c rerun (prior truths) (§25)	`36d7b07`	ignore C: `+experiment=e2c_prior_truth` C free: `+experiment=e2c_prior_truth law=three_type observation=three_split` C traj.: `+experiment=e2c_prior_truth law=three_type_traj observation=three_split` (folders `results/e2c_prior/<arm>`)
Benchmark (App. A)	`811d8ca` (gpu-bench)	`scripts/run_device_benchmark.sh`

Rows follow the order of this report. Every command has the form

```
uv run bio-sbi -m <overrides> torch_threads=4 workers=4 hydra.sweep.dir=<folder>
```

The exact command, grid and fixed settings are in each folder's `info.md`; each run's `.hydra/config.yaml` reproduces it. Packages: `bio-sbi` 0.2.0, `sbi` 0.27.0, `torch` 2.14.1+cu126, `numpy` 2.5.4, `numba` 0.68.0, `hydra-core` 1.3.8, Python 3.12.13.

Rebuilding this report (from the repository root). This report reuses the figures of the ID-ordered report (`report/figures/`, via `\graphicspath`); it has no figure script of its own:

```
uv run python report/make_figures.py      # report/figures/*.pdf, report/tables/*.csv
cd report_sequential && /home/matteo/miniconda3/envs/latex/bin/tectonic report.tex
```

`make_figures.py` only reads `results/`; incomplete sweeps would be skipped and shown as pending boxes. Numbers in tables come from the `info.md` analyses in `results/` (overview: `results/overview.md`) or from `scripts/summarize.py` functions (`mae_probabilities`, per-parameter rows); plotted numbers are in `report/tables/`.

C Glossary

A, B (C)	LGR5$^+$ stem-like, KRT20$^+$ differentiated (hidden third type)
Amortised NPE	one network trained on prior simulations, usable for any dataset
BH / GW	Bellman–Harris (continuous time, lifetimes) / Galton–Watson (generations)
C2ST	classifier two-sample test; accuracy 0.5 = indistinguishable
Clone / organoid	all descendants of one seeding event; recovered if ≥ 1 cell detected
Contraction	$1 - \mathrm{Var}_{\mathrm{post}}/\mathrm{Var}_{\mathrm{prior}}$
Control (paired)	well-specified run of the same sweep with the same seeds and datasets
Coverage (90%)	seeds whose 90% credible interval contains the truth
Design	endpoints (horizons) $\times$ organoids per endpoint
Endpoint / horizon	destructive measurement time (generations for GW, hours for BH)
Founder	seeded cell; `founder` = type mix, `founder_size` = cells per clone
Lifetime τ_A, τ_B	mean time from birth to the cell's event (BH), Gamma with shape 8
MAE	mean absolute error of posterior means over the outcome probabilities
Observed fraction	detected / seeded organoids
Paired SE	SE over seeds of the difference misspecified − control
PPC	posterior predictive check; fraction of summaries outside the 95% predictive interval
Ridge	direction in parameter space along which the data are (nearly) constant
SBC	simulation-based calibration over prior draws
SNPE	sequential NPE (rounds focus on the observed data)
Support mass	share of the NPE posterior inside the prior support
$\theta^\star$	reference truth (Table 2)
Well specified	data model = inference model